

Revisiting Underapproximate Reachability for Multipushdown Systems

S. Akshay¹, Paul Gastin², Krishna S¹, and Sparsa Roychowdhury¹

¹ IIT Bombay, Mumbai, India
{akshayss,krishnas,sparsa}@cse.iitb.ac.in
² ENS Paris-Saclay, France
paul.gastin@lsv.fr

Abstract. Boolean programs with multiple recursive threads can be captured as pushdown automata with multiple stacks. This model is Turing complete, and hence, one is often interested in analyzing a restricted class which still captures useful behaviors. In this paper, we propose a new class of bounded underapproximations for multipushdown systems, which subsumes most existing classes. We develop an efficient algorithm for solving the under-approximate reachability problem, which is based on efficient fix-point computations. We implement it in our tool **BHIM** and illustrate its applicability by generating a set of relevant benchmarks and examining its performance. As an additional takeaway **BHIM** solves the binary reachability problem in pushdown automata. To show the versatility of our approach, we then extend our algorithm to the timed setting and provide the first implementation that can handle timed multi-pushdown automata with closed guards.

Keywords: Multipushdown Systems, Underapproximate Reachability, Timed pushdown automata

1 Introduction

The reachability problem for pushdown systems with multiple stacks is known to be undecidable. However, multi-stack pushdown automata (MPDA hereafter) represent a theoretically concise and analytically useful model of multi-threaded recursive programs with shared memory. As a result, several previous works in the literature have proposed different under-approximate classes of behaviors of MPDA that can be analyzed effectively, such as *Round Bounded*, *Scope Bounded*, *Context Bounded* and *Phase Bounded* [1,2,3,4,5,6]. From a practical point of view, these underapproximations has led to efficient tools including, GetaFix [7], SPADE [8]. It has also been argued (e.g., see [9]) that such bounded underapproximations suffice to find several bugs in practice. In many such tools efficient fix-point techniques are used to speed-up computations.

We extend known fix-point based approaches by developing a new algorithm that can handle a larger class of bounded underapproximations than bounded

context and bounded scope for multi-pushdown systems while remaining efficiently implementable. This algorithm works for a new class of underapproximate behaviors called *hole bounded* behaviors, which subsumes context or scope bounded underapproximations, and is orthogonal to phase bounded underapproximations. A “hole” is a maximal sequence of push operations of a fixed stack, interspersed with well-nested sequences of any stack. Thus, in a sequence $\alpha = \beta\gamma$ where $\beta = [push_1(push_2push_3pop_3pop_2)push_1(push_3pop_3)]^{10}$ and $\gamma = push_2push_1pop_2pop_1(pop_1)^{20}$, β is a hole wrt stack 1. The suffix γ has 2 holes (the $push_2$ and the $push_1$). The number of context switches in α is > 50 , and so is the number of changes in scope, while α is 3-hole bounded. A (k)-hole bounded sequence is one such, where, at any point of the computation, the number of holes are bounded (by k). We show that the class of hole bounded sequences subsumes most of the previously defined classes of underapproximations and is, in fact, contained in the very generic class of tree-width bounded sequences. This immediately shows decidability of reachability for our class.

Analyzing the more generic class of tree-width bounded sequences is often much more difficult; for instance, building bottom-up tree automata for this purpose does not scale very well as it explores a large (and often useless) state space. Our technique is radically different from using tree automata. Under the hole-bounded assumption, we pre-compute information regarding well-nested sequences and holes using fix-point computations and use them in our algorithm. Using efficient data structures to implement this approach, we develop a tool (BHIM) for Bounded Hole reachability in Multistack pushdown systems.

Highlights of BHIM.

- Two significant aspects of the fix-point approach in BHIM are: we efficiently solve the binary reachability problem for pushdown automata. i.e., BHIM computes all pairs of states (s, t) such that t is reachable from s with empty stacks. This allows us to go beyond reachability and handle some liveness questions; (ii) we pre-compute the set of pairs of states that are endpoints of holes. This allows us to greatly limit the search for an accepting run.
- While the fix-point approach solves (binary) reachability efficiently, it does not a priori produce a witness of reachability. We remedy this situation by proposing a backtracking algorithm, which cleverly uses the computations done in the fix-point algorithm, to generate a witness efficiently.
- BHIM is parametrized w.r.t the hole bound: if non-emptiness can be checked or witnessed by a well-nested sequence (this is an easy witness and BHIM looks for easy witnesses first, then gradually increases complexity, if no easy witness is found), then it is sufficient to have the hole bound 0; increasing this complexity measure as required to certify non-emptiness gives an efficient implementation, in the sense that we search for harder witnesses only when no easier witnesses (w.r.t this complexity measure) exist. In all examples as described in the experimental section, a small (less than 4) bound suffices and we expect this to be the case for most practical examples.
- Finally, extend our approach to handle timed multi-stack pushdown systems. This shows the versatility of our approach and also requires us to solve several

technical challenges which are specific to the timed setting. Implementing this approach in **BHIM** makes it, to the best of our knowledge, the first tool that can analyze timed multi-stack pushdown automata (TMPDA) with closed guards.

We analyze the performance of **BHIM** in practice, by considering benchmarks from the literature, and generating timed variants of some of them. We modeled two variants of the Bluetooth example [10,8] and **BHIM** was able to detect three errors (of which it seems only two were already known). Likewise, for an example of a multiple producer consumer model, **BHIM** could detect bugs by finding witnesses having just 3 holes, while, it is unlikely that existing tools working on scope/context bounded underapproximations can handle them as the no. of switches in scope/context required would exceed 40 to find the bug. In the timed setting, one of the main challenges faced has been the unavailability of timed benchmarks; even in the untimed setting, many benchmarks were unavailable due to their proprietary nature. Nevertheless we tested our tool on 5 other benchmarks and 3 timed variants whose details, along with their parametric dependence plots, are given in Supplementary Material [11]. Due to lack of space proofs and technical details, especially in the timed setting are also in [11].



Related Work. Among other under-approximations, scope bounded [3] subsumes context and round bounded underapproximations, and it also paves path for GetaFix [7], a tool to analyze recursive (and multi-threaded) boolean programs. As mentioned earlier hole-boundedness strictly subsumes scope boundedness. On the other hand, GetaFix uses symbolic approaches via BDDs, which is orthogonal to the improvements made in this paper. Indeed, our next step would be to build a symbolic version of **BHIM** which extends the hole-bounded approach to work with symbolic methods. Given that **BHIM** can already handle synthetic examples with 12-13 holes (see [11]), we expect this to lead to even more drastic improvements and applicability. For sequential programs, a summary-based algorithm is used in [7]; summaries are like our well-nested sequences, except that well-nested sequences admit contexts from different stacks unlike summaries. As a result, our class of bounded hole behaviors generalizes summaries. Many other different theoretical results like phase bounded [1], order bounded [12] which gives interesting underapproximations of MPDA, are subsumed in tree-width bounded behaviors, but they do not seem to have practical implementations. Adding real-time information to pushdown automata by using clocks or timed stacks has been considered, both in the discrete and dense-timed settings. Recently, there has been a flurry of theoretical results in the topic [13,14,15,16,17]. However, to the best of our knowledge none of these algorithms have been successfully implemented (except [17] which implements a tree-automata based technique for single-stack timed systems) for multi-stack systems. One reason is that these algorithms do not employ scalable fix-point based techniques, but instead depend on region automaton-based search or tree automata-based search techniques.

2 Underapproximations in MPDA

A multi-stack pushdown automaton (MPDA) is a tuple $M = (\mathcal{S}, \Delta, s_0, \mathcal{S}_f, n, \Sigma, \Gamma)$ where, \mathcal{S} is a finite non-empty set of locations, Δ is a finite set of transitions,

127 $s_0 \in \mathcal{S}$ is the initial location, $\mathcal{S}_f \subseteq \mathcal{S}$ is a set of final locations, $n \in \mathbb{N}$ is the
 128 number of stacks, Σ is a finite input alphabet, and Γ is a finite stack alphabet
 129 which contains \perp . A transition $t \in \Delta$ can be represented as a tuple (s, op, a, s') ,
 130 where, $s, s' \in \mathcal{S}$ are respectively, the source and destination locations of the
 131 transition t , $a \in \Sigma$ is the label of the transition, and op is one of the following
 132 operations (1) **nop**, or no stack operation, (2) $(\downarrow_i \alpha)$ which pushes $\alpha \in \Gamma$ onto
 133 stack $i \in \{1, 2, \dots, n\}$, (3) $(\uparrow_i \alpha)$ which pops stack i if the top of stack i is $\alpha \in \Gamma$.

134 For a transition $t = (s, \text{op}, a, s')$ we write $\text{src}(t) = s$, $\text{tgt}(t) = s'$ and $\text{op}(t) = \text{op}$.
 135 At the moment we ignore the action label a but this will be useful later when we
 136 go beyond reachability to model checking. A *configuration* of the MPDA is a tuple
 137 $(s, \lambda_1, \lambda_2, \dots, \lambda_n)$ such that, $s \in \mathcal{S}$ is the current location and $\lambda_i \in \Gamma^*$ represents
 138 the current content of i^{th} stack. The semantics of the MPDA is defined as follows:
 139 a run is accepting if it starts from the initial state and reaches a final state with
 140 all stacks empty. The language accepted by a MPDA is defined as the set of words
 141 generated by the accepting runs of the MPDA. Since the reachability problem for
 142 MPDA is Turing complete, we consider under-approximate reachability.

143 A sequence of transitions is called **complete** if each push in that sequence
 144 has a matching pop and vice versa. A **well-nested** sequence denoted ws is
 145 defined inductively as follows: a possibly empty sequence of **nop**-transitions is
 146 ws , and so is the sequence $t ws t'$ where $\text{op}(t) = (\downarrow_i \alpha)$ and $\text{op}(t') = (\uparrow_i \alpha)$ are a
 147 matching pair of push and pop operations of stack i . Finally the concatenation
 148 of two well-nested sequences is a well-nested sequence, i.e., they are closed under
 149 concatenation. The set of all well-nested sequences defined by an MPDA is
 150 denoted **WS**. If we visualize this by drawing edges between pushes and their
 151 corresponding pops, well-nested sequences have no crossing edges, as in 
 152 and , where we have two stacks, depicted with red and violet edges. We
 153 emphasize that a well-nested sequence can have well-nested edges from any stack.
 154 In a sequence σ , a push (pop) is called a **pending** push (pop) if its matching
 155 pop (push) is not in the same sequence σ .

156 **Bounded Underapproximations.** As mentioned in the introduction, different
 157 bounded under-approximations have been considered in the literature to get
 158 around the Turing completeness of MPDA. During a computation, a context is a
 159 sequence of transitions where only one stack or no stack is used. In *context bounded*
 160 computations the number of contexts are bounded [18]. A *round* is a sequence
 161 of (possibly empty) contexts for stacks $1, 2, \dots, n$. *Round bounded* computations
 162 restrict the total number of rounds allowed [2,16,17]. *Scope bounded* computations
 163 generalize bounded context computations. Here, the context changes within any
 164 push and its corresponding pop is bounded [2,5,6]. A *phase* is a contiguous
 165 sequence of transitions in a computation, where we restrict pop to only one stack,
 166 but there are no restrictions on the pushes [1]. A phase bounded computation is
 167 one where the number of phase changes is bounded.

168 **Tree-width.** A generic way of looking at them is to consider classes which have a
 169 bound on the tree-width [19]. In fact, the notions of split-width/cliue-width/tree-
 170 width of communicating finite state machines/timed push down systems has been
 171 explored in [20], [21]. The behaviors of the underlying system are then represented

as graphs. It has been shown in these references that if the family of graphs arising from the behaviours of the underlying system (say S) have a bounded tree-width, then the reachability problem is decidable for S via, tree-automata. However, this does not immediately give rise to an efficient implementation. The tree-automata approach usually gives non-deterministic or bottom-up tree automata, which when implemented in practice (see [17]) tend to blow up in size and explore a large and useless space. Hence there is a need for efficient algorithms, which exist for more specific underapproximations such as context-bounded (leading to fix-point algorithms and their practical implementations [7]).

2.1 A new class of under-approximations

Our goal is to bridge the gap between having practically efficient algorithms and handling more expressive classes of under-approximations for reachability of multi-stack pushdown systems. To do so, we define a bounded approximation which is expressive enough to cover previously defined practically interesting classes (such as context bounded etc), while at the same time allowing efficient decidable reachability tests, as we will see in the next section.

Definition 1. (*Holes*). Let σ be complete sequence of transitions, of length n in a MPDA, and let ws be a (possibly empty) well-nested sequence.

- A **hole** of stack i is a maximal factor of σ of the form $(\downarrow_i ws)^+$, where $ws \in \text{WS}$. The maximality of the hole of stack i follows from the fact that any possible extension ceases to be a hole of stack i ; that is, the only possible events following a maximal hole of stack i are a push \downarrow_j of some stack $j \neq i$, or a pop of some stack $j \neq i$. In general, whenever we speak about a hole, the underlying stack is clear.
- A push \downarrow_i in a hole (of stack i) is called a pending push at (i.e., just before) a position $x \leq n$, if its matching pop occurs in σ at a position $z > x$.
- A hole (of stack i) is said to be **open** at a position $x \leq n$, if there is a pending push \downarrow_i of the hole at x . Let $\#_x(\text{hole})$ denote the number of open holes at position x . The **hole bound** of σ is defined as $\max_{1 \leq x \leq |\sigma|} \#_x(\text{hole})$.
- A hole segment of stack i is a prefix of a hole of stack i , ending in a ws , while an atomic hole segment of stack i is just the segment of the form $\downarrow_i ws$.

As an example, consider the sequence σ in Figure 1 of transitions of a MPDA having stacks 1,2 (denoted respectively red and blue). We use superscripts for each push, pop of each stack to distinguish the i th push, j th pop and so on of each stack. There are two holes of stack 1 (red stack) denoted by the red patches,




Fig. 1. A run σ with 2 holes (2 red patches) of the red stack and 1 hole (one blue patch) of the blue stack.

and one hole of stack 2 (blue stack) denoted by the blue patch. The subsequence $\downarrow_1^1 \downarrow_1^2 ws_2$ of the first hole is not a maximal factor, since it can be extended by $\downarrow_1^3 ws_3$ in the run σ , extending the hole. Consider the position in σ marked with \downarrow_2^1 . At this position, there is an open hole of the red stack (the first red patch), and there is an open hole of the blue stack (the blue patch). Likewise, at the position \uparrow_1^5 , there are 2 open holes of the red stack (2 red patches) and one open hole of the blue stack 2 (the blue patch). The hole bound of σ is 3. The green patch consisting of $\uparrow_1^3, \uparrow_1^2$ and ws_5 is a pop-hole of stack 1. Likewise, the pops $\uparrow_2^2, \uparrow_1^5, \uparrow_2^1$ are all pop-holes (of length 1) of stacks 2,1,2 respectively.

Definition 2. (HOLE BOUNDED REACHABILITY PROBLEM) *Given a MPDA and $K \in \mathbb{N}$, the K -hole bounded reachability problem is the following: Does there exist a K -hole bounded accepting run of the MPDA?*

Proposition 1. *The tree-width of K -hole bounded MPDA behaviors is at most $(2K + 3)$.*

A detailed proof of this Proposition is given in Appendix A.1. Once we have this, from [19][16][17], decidability and complexity follow immediately. Thus,

Corollary 1. *The K -hole bounded reachability problem for MPDA is decidable in $\mathcal{O}(|\mathcal{M}|^{2K+3})$ where, \mathcal{M} is the size of the underlying MPDA.*

Next, we turn to the expressiveness of this class wrt to the classical underapproximations of MPDA: first, the **hole** bounded class strictly subsumes **scope** bounded which already subsumes **context** bounded and **round** bounded classes. Also **hole** bounded MPDA and **phase** bounded MPDA are orthogonal.

Proposition 2. *Consider a MPDA M . For any K , let L_K denote a set of sequences accepted by M which have number of rounds or number of contexts or scope bounded by K . Then there exists $K' \leq K$ such that L_K is K' hole bounded. Moreover, there exist languages which are K hole bounded for some constant K , which are not K' round or context or scope bounded for any K' . Finally, there exists a language which is accepted by phase bounded MPDA but not accepted by hole bounded MPDA and vice versa.*

Proof. We first recall that if a language L is K -round, or K -context bounded, then it is also K' -scope bounded for some $K' \leq K$ [5,2]. Hence, we only show that scope bounded systems are subsumed by hole bounded systems.

Let L be a K -scope bounded language, and let M be a MPDA accepting L . Consider a run ρ of $w \in L$ in M . Assume that at any point i in the run ρ , $\#_i(\text{holes}) = k'$, and towards a contradiction, let, $k' > K$. Consider the leftmost open hole in ρ which has a pending push \downarrow^p whose pop \uparrow^p is to the right of i . Since $k' > K$ is the number of open holes at i , there are at least $k' > K$ context changes in between \downarrow^p and \uparrow^p . This contradicts the K -scope bounded assumption, and hence $k' \leq K$.

To show the strict containment, consider the visibly pushdown language [22] given by $L^{bh} = \{a^n b^n (a^{p_1} c^{p_1+1} b^{p'_1} d^{p'_1+1} \dots a^{p_n} c^{p_n+1} b^{p'_n} d^{p'_n+1}) \mid n, p_1, p'_1, \dots, p_n, p'_n \in$

248 $\mathbb{N}\}$. A possible word $w \in L^{bh}$ is $a^3b^3 a^2c^3b^2d^3 a^2c^3bd^2 ac^2bd^2$ with a, b representing
 249 push in stack 1,2 respectively and c, d representing the corresponding matching
 250 pop from stack 1,2. A run ρ accepting the word $w \in L^{bh}$ will start with a sequence
 251 of pushes of stack 1 followed by another sequence of pushes of stack 2. Note that,
 252 the number of the pushes n is same in both stacks. Then there is a group G
 253 consisting of a well-nested sequence of stack 1 (equal a and c) followed by a pop
 254 of the stack 1 (an extra c), another well-nested sequence of stack 2 (equal b and
 255 d) and a pop of the stack 2 (an extra d), repeated n times. From the definition
 256 of the `hole`, the total number of holes required in G is 0. But, we need 1 hole for
 257 the sequence of a 's and another for the sequence of b 's at the beginning of the
 258 run, which creates at most 2 holes during the run. Thus, the hole bound for any
 259 accepting run ρ is 2, and the language L^{bh} is 2-hole bounded.

However, L^{bh} is not k -scope bounded for any k . Indeed, for each $m \geq 1$,
 consider the word $w_m = a^mb^m(ac^2bd^2)^m \in L^{bh}$. It is easy to see that w_m is
 2m-scope bounded (the matching c, d of each a, b happens $2m$ context switches
 later) but not k -scope bounded for $k < 2m$. It can be seen that L^{bh} is not k -phase
 bounded either. Finally, $L' = \{(ab)^nc^nd^n \mid n \in \mathbb{N}\}$ with a, b and c, d respectively
 being push and pop of stack 1,2 is not hole-bounded but 2-phase bounded. \square

260 3 A Fix-point Algorithm for Hole Bounded Reachability

261 In the previous section, we showed that hole-bounded underapproximations are
 262 a decidable subclass for reachability, by showing that this class has a bounded
 263 tree-width. However, as explained in the introduction, this does not immediately
 264 give a fix-point based algorithm, which has been shown to be much more efficient
 265 for other more restricted sub-classes, e.g., context-bounded. In this section, we
 266 provide such a fix-point based algorithm for the hole-bounded class and explain its
 267 advantages. Later we discuss its versatility by showing extensions and evaluating
 268 its performance on a suite of benchmarks.

269 We describe the algorithm in two steps: first we give a simple fix-point based
 270 algorithm for the problem of 0-hole or *well-nested reachability*, i.e, reachability
 271 by a well-nested sequence without any holes. For the 0-hole case, our algorithm
 272 computes the *reachability relation*, also called the *binary reachability problem* [23].
 273 That is, we accept all pairs of states (s, s') such that there is a well-nested run
 274 from s with empty stack to s' with empty stack. Subsequently, we combine this
 275 binary reachability for well-nested sequences with an efficient graph search to
 276 obtain an algorithm for K -hole bounded reachability.

277 **Binary well-nested reachability for MPDA.** Note that single stack PDA are
 278 a special case, since all runs are indeed well-nested.

- 279 1. **Transitive Closure:** Let \mathcal{R} be the set of tuples of the form (s_i, s_j) representing
 280 that state s_j is reachable from state s_i via a `nop` discrete transition. Such a
 281 sequence from s_i to s_j is trivially *well-nested*. We take the **TransitiveClosure**
 282 of \mathcal{R} using Floyd-Warshall algorithm [24]. The resulting set \mathcal{R}_c of tuples
 283 answers the binary reachability for finite state automata (no stacks).

Algorithm 1: Algorithm for Emptiness Checking of hole bounded MPDA

```

1 Function IsEmpty( $M = (\mathcal{S}, \Delta, s_0, \mathcal{S}_f, n, \Sigma, \Gamma), K$ ):
   Result: True or False
2  $WR := \text{WellNestedReach}(M);$  \\Solves binary reachability for pushdown system
3 if some  $(s_0, s_1) \in WR$  with  $s_1 \in \mathcal{S}_f$  then
4   return False;
5 forall  $i \in [n]$  do
6    $AHS_i := \emptyset; Set_i := \emptyset;$ 
7   forall  $(s, \downarrow_i(\alpha), a, s_1) \in \Delta$  and  $(s_1, s') \in WR$  do
8      $AHS_i := AHS_i \cup \{(i, s, \alpha, s')\}; Set_i := Set_i \cup \{(s, s')\};$ 
9      $HS_i := \{(i, s, s') \mid (s, s') \in \text{TransitiveClosure}(Set_i)\};$ 
10   $\mu := [s_0]; \mu.\text{NumberOfHoles} := 0;$ 
11   $SetOfLists_{new} := \{\mu\}; SetOfLists := \emptyset;$ 
12  do
13     $SetOfLists := SetOfLists \cup SetOfLists_{new};$ 
14     $SetOfLists_{todo} := SetOfLists_{new}; SetOfLists_{new} := \emptyset;$ 
15    forall  $\mu' \in SetOfLists_{todo}$  do
16      if  $\mu'.\text{NumberOfHoles} < K$  then
17        forall  $i \in [n]$  do
18          \\ Add hole for stack i
19           $SetOfLists_h := \text{AddHole}_i(\mu', HS_i) \setminus SetOfLists;$ 
20           $SetOfLists_{new} := SetOfLists_{new} \cup SetOfLists_h;$ 
21        if  $\mu'.\text{NumberOfHoles} > 0$  then
22          forall  $i \in [n]$  do
23            \\ Add pop for stack i
24             $SetOfLists_p := \text{AddPop}_i(\mu', M, AHS_i, HS_i, WR) \setminus SetOfLists;$ 
25             $SetOfLists_{new} := SetOfLists_{new} \cup SetOfLists_p;$ 
26            forall  $\mu_3 \in SetOfLists_p$  do
27              if  $\mu_3.\text{last} \in \mathcal{S}_f$  and  $\mu_3.\text{NumberOfHoles} = 0$  then
28                return False; \\If reached destination state
29  while  $SetOfLists_{new} \neq \emptyset;$ 
30  return True;

```

284 **2. Push-Pop Closure:** For stack operations, consider a push transition on
 285 some stack (say stack i) of symbol γ , enabled from a state s_1 , reaching state
 286 s_2 . If there is a matching pop transition from a state s_3 to s_4 , which pops the
 287 same stack symbol γ from the stack i and if we have $(s_2, s_3) \in \mathcal{R}_c$, then we
 288 can add the tuple (s_1, s_4) to \mathcal{R}_c . The function `WellNestedReach` (Algorithm 2,
 289 Appendix B) repeats this process and the transitive closure described above
 290 until a fix-point is reached. Let us denote the resulting set of tuples by WR .
 291 Thus, we have

292 **Lemma 1.** $(s_1, s_2) \in WR$ iff \exists a well-nested run in the MPDA from s_1 to s_2 .

293 **Beyond well-nested reachability.** A naive algorithm for K -hole bounded
 294 reachability for $K > 0$ is to start from the initial state s_0 , and do a Breadth
 295 First Search (BFS), nondeterministically choosing between extending with a
 296 well-nested segment, creating hole segments (with a pending push) and closing
 297 hole segments (using pops). We accept when there are no open hole segments
 298 and reach a final state; this gives an exponential time algorithm. Given the
 299 exponential dependence on the hole-bound K (Corollary 1), this exponential
 300 blowup is unavoidable in the worst case, but we can do much better in practice.
 301 In particular, the naive algorithm makes arbitrary non-deterministic choices
 302 resulting in a blind exploration of the BFS tree.

303 In this section, we use the binary well-nested reachability algorithm as an
 304 efficient subroutine to limit the search in BFS to its reachable part (note that
 305 this is quite different from DFS as well since we do not just go down one path).
 306 The crux is that at any point, we create a new hole for stack i , *only* when (i)
 307 we know that we cannot reach the final state without creating this hole and (ii)
 308 we know that we can close all such holes which have been created. Checking (i)
 309 is easy, since we just use the WR relation for this. Checking (ii) blindly would
 310 correspond to doing a DFS; however, we precompute this information and simply
 311 look it up, resulting in a constant time operation after the precomputation.

312 **Precomputing hole information.** Recall that a *hole* of stack i is a maximal
 313 sequence of the form $(\downarrow_i ws)^+$, where ws is a well-nested sequence and \downarrow_i
 314 represents a push of stack i . A *hole segment* of stack i is a prefix of a hole
 315 of stack i , ending in a ws , while an *atomic hole segment* of stack i is just the
 316 segment of the form $\downarrow_i ws$. A *hole-segment* of stack i which starts from state s
 317 in the MPDA and ends in state s' , can be represented by the triple (i, s, s') , that
 318 we call a *hole triple*. We compute the set HS_i of all hole triples (i, s, s') such that
 319 starting at s , there is a hole segment of stack i which ends at state s' , as detailed
 320 in lines (5-9) of Algorithm 1. In doing so, we also compute the set AHS_i of all
 321 atomic hole segments of stack i and store them as tuples of the form (i, s_p, α, s_q)
 322 such that s_p and s_q are the MPDA states respectively at the left and right end
 323 points of an atomic hole segment of stack i , and α is the symbol pushed on stack
 324 i ($s_p \xrightarrow{\downarrow_i(\alpha)ws} s_q$).

325 **A guided BFS exploration.** We start with a list $\mu_0 = [s_0]$ consisting of
 326 the initial state and construct a BFS exploration tree whose nodes are lists of
 327 bounded length. A list is a sequence of states and hole triples representing a
 328 K -hole bounded run in a concise form. If H_i represents a hole triple for stack i ,
 329 then a list is a sequence of the form $[s, H_i, H_j, H_k, H_i, \dots, H_\ell, s']$. The simplest
 330 kind of list is a single state s . For example, a list with 3 holes of stacks i, j, k is
 331 $\mu = [s_0, (i, s, s'), (j, r, r'), (k, t, t'), t']$. The hole triples (in red) denote open holes
 332 in the list. The maximum number of open holes in a list is bounded, making the
 333 length of the list also bounded. Let $\text{last}(\mu)$ represent the last element of the list
 334 μ . This is always a state. For a node v storing list μ in the BFS tree, if v_1, \dots, v_k
 335 are its children, then the corresponding lists μ_1, \dots, μ_k are obtained by extending
 336 the list μ by one of the following operations:

- 337 1. **Extend μ with a hole.** Assume there is a hole of some stack i , which starts
 338 at $\text{last}(\mu) = s$, and ends at s' . If the list at the parent node v is $\mu = [\dots, s]$,
 339 then for all $(i, s, s') \in HS_i$, we obtain the list $\text{trunc}(\mu) \cdot \text{append}[(i, s, s'), s']$
 340 at the child node (i.e., we remove the last element s of μ , then append to
 341 this list the hole triple (i, s, s') , followed by s'). Algorithm 3 in Appendix
 342 describes this operation in more detail.
- 343 2. **Extend μ with a pop.** Suppose there is a transition $t = (s_k, \uparrow_i(\alpha), a, s'_k)$
 344 from $\text{last}(\mu) = s_k$, where μ is of the form $[s_0, \dots, (h, u, v), (i, s, s'), (j, t, t'), \dots, s_k]$,
 345 such that there is no hole triple of stack i after (i, s, s') , we extend the run by
 346 matching this pop (with its push). However, to obtain the last pending push

of stack i corresponding to this hole, just HS_i information is not enough since we also need to match the stack content. Instead, we check if we can split the hole (i, s, s') into (1) a hole triple $(i, s, s_a) \in HS_i$, and (2) a tuple $(i, s_a, \alpha, s') \in AHS_i$. If both (1) and (2) are possible, then the pop transition t corresponds to the last pending push of the hole (i, s, s') . t indeed matches the pending push recorded in the atomic hole (i, s_a, α, s') in μ , enabling the firing of transition t from the state s_k , reaching s'_k . In this case, we add the child node with the list μ' obtained from μ as follows. We replace (i) s_k with s'_k , and (ii) (i, s, s') with (i, s, s_a) , respectively signifying firing of the transition t and the “shrinking” of the hole, by shifting the end point of the hole segment to the left. When we obtain the hole triple (i, s, s) (the start and end points of the hole segment coincide), we may have uncovered the last pending push and thereby “closed” the hole segment completely. At this point, we may choose to remove (i, s, s) from the list, obtaining $[s_0, \dots, (h, u, v), (j, t, t') \dots, s'_k]$. For every such $\mu' = [s_0, \dots, (h, u, v), (i, s, s_a), (j, t, t'), \dots, s'_k]$ and all $(s'_k, s_m) \in WS$ we also extend μ' to $\mu'' = [s_0, \dots, (h, u, v), (i, s, s_a), (j, t, t'), \dots, s_m]$. Notice that the size of the list in the child node obtained on a pop, is either the same as the list in the parent, or is smaller. The details are in Algorithm 4.

The number of lists is bounded since the number of states and the length of the lists are bounded. The BFS exploration tree will thus terminate. Combining the above steps gives us Algorithm 1, whose correctness gives us:

Theorem 1. *Given a MPDA and a positive integer K , Algorithm 1 always terminates and answers “false” iff there exists a K -hole bounded accepting run of the MPDA.*

Complexity of the Algorithm. The maximum number of states of the system is $|\mathcal{S}|$. The time complexity of transitive closure is $\mathcal{O}(|\mathcal{S}|^3)$, using a Floyd-Warshall implementation. The time complexity of Algorithm 2, which uses the transitive closure, is $\mathcal{O}(|\mathcal{S}|^5) + \mathcal{O}(|\mathcal{S}|^2 \times (|\Delta| \times |\mathcal{S}|))$. To compute AHS for n stacks the time complexity is $\mathcal{O}(n \times |\Delta| \times |\mathcal{S}|^2)$ and to compute HS for n stacks the complexity is $\mathcal{O}(n \times |\mathcal{S}|^2)$. For multistack systems, each list keeps track of (i) the number of hole segments ($\leq K$), and (ii) information pertaining to holes (start, end points of holes, and which stack the hole corresponds to). In the worst case, this will be $(2K + 2)$ possible states in a list, as we are keeping the states at the start and end points of all the hole segments and a stack per hole. So, there are $\leq |\mathcal{S}|^{2K+3} \times n^{K+1}$ lists. In the worst case, when there is no K -hole bounded run, we may end up generating all possible lists for a given bound K on the hole segments. The time complexity is thus bounded above by $\mathcal{O}(|\mathcal{S}|^{2K+3} \times n^{K+1} + |\mathcal{S}|^5 + |\mathcal{S}|^3 \times |\Delta|)$.

Beyond Reachability. We can solve the usual safety questions in the (bounded-hole) underapproximate setting, by checking for underapproximate reachability on the product of the given system with the complement of the safe set. Given the way Algorithm 1 is designed, the fix-point algorithm allows us to go beyond reachability. In particular, we can solve several (increasingly difficult) variants of the repeated reachability problem, without much modification.

Consider the question : For a given state s and MPDA, does there exist a run ρ starting from s_0 which visits s infinitely often? This is decidable if we can

392 decompose ρ into a finite prefix ρ_1 and an infinite suffix ρ_2 s.t. (1) Both ρ_1, ρ_2
 393 are well-nested, or (2) ρ_1 is K -hole bounded complete (all stacks empty), and ρ_2
 394 is well-nested, or (3) ρ_1 is K -hole bounded, and $\rho_2 = (\rho_3)^\omega$, where ρ_3 is K -hole
 395 bounded. It is easy to see that (1) is solved by two calls to `WellNestedReach` and
 396 choosing non-empty runs. (2) is solved by a call to Algorithm 1, modified so that
 397 we reach s , and then calling `WellNestedReach`. Lastly, to solve (3), first modify
 398 Algorithm 1 to check reachability to s with possibly non-empty stacks. Then run
 399 the modified algorithm twice : first start from s_0 and reach s ; second start from
 400 s and reach s again.

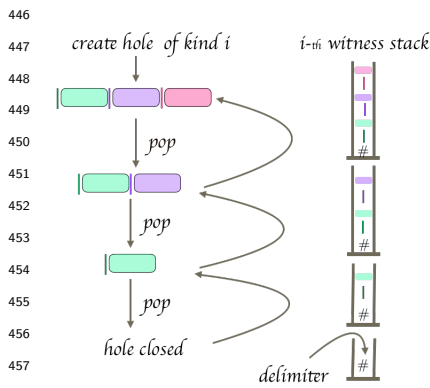
401 4 Generating a Witness

402 We next focus on the question of generating a witness for an accepting run
 403 when our algorithm guarantees non-emptiness. This question is important to
 404 address from the point of view of applicability: if our goal is to see if bad states
 405 are reachable, i.e., non-emptiness corresponds to presence of a bug, the witness
 406 run gives the trace of how the bug came about and hence points to what can
 407 be done to fix it (e.g., designing a controller). We remark that this question is
 408 difficult in general. While there are naive algorithms which can explore for the
 409 witness (thus also solving reachability), these do not use fix-point techniques and
 410 hence are not efficient. On the other hand, since we use fix-point computations
 411 to speed up our reachability algorithm, finding a witness, i.e., an explicit run
 412 witnessing reachability, becomes non-trivial. Generation of a witness in the case of
 413 well-nested runs is simpler than the case when the run has holes, and requires us
 414 to “unroll” pairs $(s_0, s_f) \in \mathbf{WR}$ recursively and generate the sequence of transitions
 415 responsible for (s_0, s_f) , as detailed in Algorithm 5.

416 **Getting Witnesses from Holes.** Now we move on to the more complicated
 417 case of behaviours having holes. Recall that in BFS exploration we start from
 418 the states reachable from s_0 by well-nested sequences, and explore subsequent
 419 states obtained either from (i) a hole creation, or (ii) a pop operation on a stack.
 420 Proceeding in this manner, if we reach a final configuration (say s_f), with all
 421 holes closed (which implies empty stacks), then we declare non-emptiness. To
 422 generate a witness, we start from the final state s_f reachable in the run (a leaf
 423 node in the BFS exploration tree) and *backtrack* on the BFS exploration tree
 424 till we reach the initial state s_0 . This results in generating a witness run in the
 425 reverse, from the right to the left.

426 • Assume that the current node of the BFS tree was obtained using a pop
 427 operation. There are two possibilities to consider here (see below) depending on
 428 whether this pop operation closed or shrunk some hole. Recall that each hole
 429 has a left end point and a right end point and is of a specific stack i , depending
 430 on the pending pushes \downarrow_i it has. So, if the MPDA has k stacks, then a list in the
 431 exploration tree can have k kinds of holes. The witness algorithm uses k stacks
 432 called *witness stacks* to correctly implement the backtracking procedure, to deal
 433 with k kinds of holes. Witness stacks should not be confused with the stacks of
 434 the MPDA.

435 • Assume that the current pop operation is closing a hole $| \text{ } \text{ } \text{ } |$ of kind
 436 i as in Figure 2. This hole consists of the atomic holes $| \text{ } \text{ } |$, $| \text{ } \text{ } |$ and $| \text{ } \text{ } |$. The
 437 atomic hole $| \text{ } \text{ } |$ consists of the push $|$ and the well-nested sequence $| \text{ } \text{ } |$ (same
 438 for the other two atomic holes). Searching among possible push transitions, we
 439 identify the matching push $|$ associated with the current pop, resulting in closing
 440 the hole. On backtracking, this leads to a parent node with the atomic hole $| \text{ } \text{ } |$
 441 having as left end point, the push $|$, and the right end point as the target of
 442 the ws $| \text{ } \text{ } |$. We push onto the witness stack i , a barrier (a delimiter symbol $\#$)
 443 followed by the matching push transition $|$ and then the ws , $| \text{ } \text{ } |$. The barrier
 444 segregates the contents of the witness stack when we have two pop transitions of
 445 the same stack in the reverse run, closing/shrinking two different holes.



459 **Fig. 2.** Backtracking to spit out
 460 the hole $| \text{ } \text{ } \text{ } |$ in reverse. The
 461 transitions of the atomic hole $| \text{ } \text{ } |$
 462 are first written in the reverse
 463 order, followed by those of $| \text{ } \text{ } |$ in
 464 reverse, and then of $| \text{ } \text{ } |$ in reverse.

465 hole segments which constituted this hole. Notice that when we finish processing
 466 a hole of kind i , then the witness stack i has the hole reversed inside it, followed
 467 by a barrier. The next hole of the same kind i will be treated in the same manner.

468 • If the current node of the BFS tree is obtained by creating a hole of kind i
 469 in the fix-point algorithm, then we pop the contents of witness stack i till we
 470 reach a barrier. This spits out the atomic hole segments of the hole from the
 471 right to the left, giving us a sequence of push transitions, and the respective ws
 472 in between. The transitions constituting the ws are retrieved using Algorithm 5
 473 and added. Notice that popping the witness stack i till a barrier spits out the
 474 sequence of transitions in the correct reverse order while backtracking.

475 5 Adding Time to Multi-pushdown systems

476 In this section, we briefly describe how the algorithms described in section 3
 477 can be extended to work in the timed setting. Due to lack of space, we focus

478 on some of the significant challenges and advances, leaving the formal details
 479 and algorithms to the supplement [11]. A TMPDA extends a MPDA with clock
 480 variables. Transitions check constraints which are conjunctions/disjunctions of
 481 constraints (called closed guards in the literature) of the form $x \leq c$ or $x \geq c$ for
 482 $c \in \mathbb{N}$ and x any clock. Symbols pushed on stacks “age” with time elapse. A pop
 483 is successful only when the age of the symbol lies within a certain interval. The
 484 acceptance condition is as in the case of MPDA.

485 The first main challenge in adapting the algorithms in section 3 to the timed
 486 setting was to take care of all possible time elapses along with the operations
 487 defined in Algorithm 1. The usage of closed guards in TMPDA means that it
 488 suffices to explore all runs with integral time elapses (for a proof see e.g., Lemma
 489 4.1 in [16]). Thus configurations are pairs of states with valuations that are vectors
 490 of non-negative integers, each of which is bounded by the maximal constant in
 491 the system. Now, to check reachability we need to extend all the precomputations
 492 (transitive closure, well-nested reachability, as well as atomic and non-atomic hole
 493 segments) with the time elapse information. To do this, we use a weighted version
 494 of the Floyd-Warshall algorithm by storing time elapses during precomputations.
 495 This allows us to use this precomputed *timed* well-nested reachability information
 496 while performing the BFS tree exploration, thus ensuring that any explored state
 497 is indeed reachable by a timed run. In doing so, the most challenging part is
 498 extending the BFS tree wrt a pop. Here, we not only have to find a split of a
 499 hole into an atomic hole-segment and a hole-segment as in Algorithm 1, but also
 500 need to keep track of possible partitions of time.

501 **Timed Witness:** As in the untimed case, we generate a witness certifying non-
 502 emptiness of TMPDA. But, producing a witness for the fix-point computation
 503 as discussed earlier requires unrolling. The fix-point computation generates a
 504 pre-computed set **WRT** of tuples $((s, \nu), t, (s', \nu'))$, where $s, s' \in \mathcal{S}$, t is time elapsed
 505 in the well-nested sequence and $\nu, \nu' \in \mathbb{N}^{|\mathcal{X}|}$ are integral valuations. This set
 506 of tuples does not have information about the intermediate transitions and
 507 time-elapses. To handle this, using the pre-computed information, we define a
 508 lexicographic progress measure which ensures termination of this search.

509 While the details are in [11] (Algorithm 14), the main idea is as follows:
 510 the first progress measure is to check if there a time-elapse t transition possible
 511 between (s, ν) and (s', ν') and if so, we print this out. If not, $\nu' \neq \nu + t$, and
 512 some set of clocks have been reset in the transition(s) from (s, ν) to (s', ν') . The
 513 second progress measure looks at the sequence of transitions from (s, ν) to (s', ν') ,
 514 consisting of reset transitions (at most the number of clocks) that result in ν'
 515 from ν . If neither the first nor the second progress measure apply, then $\nu = \nu'$,
 516 and we are left to explore the last progress measure, by exploring at most $|\mathcal{S}|$
 517 number of transitions from (s, ν) to (s', ν') . The lexicographic progress measure
 518 seamlessly extends the witness generation to the timed setting.

519 6 Implementation and Experiments

520 We implemented a tool BHIM (**B**ounded **H**oles **I**n **M**PD**A**) written in C++ based
521 on Algorithm 1, which takes an MPDA and a constant K as input and returns
522 (*True*) iff there exists a K -hole bounded run from the start state to an accepting
523 state of the MPDA. In case there is such an accepting run, BHIM generates one
524 such, with minimal number of holes. For a given hole bound K , BHIM first tries
525 to produce a witness with 0 holes, and iteratively tries to obtain a witness by
526 increasing the bound on holes till K . In most of the cases, BHIM found the
527 witness before reaching the bound K . Whenever BHIM's witness had K holes, it
528 is guaranteed that there are no witnesses with a smaller number of holes.

529 To evaluate the performance of BHIM, we looked at some available benchmarks
530 and modeled them as MPDA. We also added timing constraints to some examples
531 such that they can be modeled as TMPDA. Our tests were run on a GNU/Linux
532 system with Intel® Core™ i7-4770K CPU @ 3.50GHz, and 16GB of RAM. We
533 considered overall 7 benchmarks, of which we sketch 3 in detail here. The details
534 of these as well as the remaining ones are in [11].

535 • **Bluetooth Driver [18]**. The Bluetooth device driver example [18], has two
536 threads and a shared memory. We model this driver using a 2-stack pushdown
537 system, where a state represents the current valuation of the global variables and
538 stacks are used to maintain the call-return between different functions and to
539 keep the count of processes currently using the driver. There is also a scheduler
540 which can preempt any thread executing a non-atomic instruction. A known error
541 as pointed out in [18] is a race condition between two threads where one thread
542 tries to write to a global variable and the other thread tries to read from it. BHIM
543 found this error, with a well-nested witness. A timed extension of this example
544 was also considered, where, a witness was obtained again with hole bound 0.

545 • **A Multi-threaded Producer Consumer Problem**. The Producer consumer
546 problem (see e.g., [25]) is a classic example of concurrency and synchronization.
547 A more challenging version of this is when there are multiple producers and
548 consumers. Assume that two ingredients called 'A' and 'B' are produced in a
549 production line in batches, where a batch can produce arbitrarily many items,
550 but it fixed for a day. Further, assume that (1) two units of 'A' and one unit of
551 'B' make an item called 'C'; (2) the production line begins its day by producing
552 a batch of A's and then the rest of the day, it keeps producing B's in batches,
553 one after the other. During the day, 'C's are churned out using 'A' and 'B' in the
554 proportion mentioned above and, if we run out of 'A's, we obtain an error; there
555 is no problem if 'B' is exhausted, since a fresh batch producing 'B' is commenced.
556 This idea can be imagined as a real life scenario where item 'A' represents an
557 item which is very expensive to produce but can be produced in large amount
558 but the item 'B' can be produced frequently, but it has to be consumed very
559 soon, if it is not consumed then it becomes useless.

560 For $p, q \in \mathbb{N}$, consider words of the form $a^m(b^k(c^2d)^k)^n$ where, a represents
561 the production of one unit of 'A', b represents the production of one unit of 'B',
562 c represents consumption of one unit of 'A' and d represents consumption of
563 one unit of 'B'. 'm' represents the production capacity of 'A' for the day and 'k'

564 represents production capacity of ‘B’(per batch) for the day, ‘n’ represents the
565 number batches of ‘B’ produced in a day. Unless $m \geq 2nk$, we will obtain an
566 error. This is easily modeled using a 2 stack visibly multi pushdown automaton
567 where a, b are push symbols of stack 1,2 respectively and c, d are pop symbols
568 of stack 1,2 respectively. Let $L_{m,k,n}$ be the set of words of the above form s.t.
569 $2nk < m$. It can be seen that $L_{m,k,n}$ does not have any well-nested word in it.
570 The number of context switches(also, scope bound) in words of $L_{m,k,n}$ depends
571 on the parameters k and n . However, $L_{m,k,n}$ is 2 hole-bounded : at any position
572 of the word, the open holes come from the unmatched sequences of a and b seen
573 so far. BHIM checked for the non-emptiness of $L_{m,k,n}$ with a witness of hole
574 bound 2.

575 • **Critical time constraints [26]**. This is one of the timed examples, where
576 we consider the language $L^{crit} = \{a^y b^z c^y d^z \mid y, z \geq 1\}$ with time constraints
577 between occurrences of symbols. The first c must appear after 1 time-unit of the
578 last a , the first d must appear within 3 time-units of the last b , and the last b
579 must appear within 2 time units from the start, and the last d must appear at 4
580 time units. L^{crit} is accepted by a TMPDA with two timed stacks. L^{crit} has no
581 well-nested word, is 4-context bounded, but only 2 hole-bounded. Note that this
582 example can also be modeled using 2 un-timed stacks with 4 clocks to maintain
583 the timing constraints. Clearly, the model with timed stack gives us a more
584 succinct representation which also avoid the state space explosion due to added
585 clocks. In experiments, the added clocks increased the run time exponentially,
586 which clearly shows the advantage of modeling using timed stack. The detail
587 comparison of run time between this two models is shown in the Table 2. The
588 “age” column indicates weather the stack is timed or not.

589 • **A Linux Kernel bug dm_target.c [27]**. This example is about a double
590 free bug in the file `drivers/md/dm-target.c` in Linux Kernel 2.5.71, which
591 was introduced to fix a memory leak, but it ended up double freeing the object.
592 BHIM found this bug with a witness of hole bound 3.

593 • **Concurrent Insertions in Binary Search Trees**. Concurrent insertions in
594 binary search trees is a very important problem in database management systems.
595 [28] proposes an algorithm to solve this problem for concurrent implementations.
596 However, if the locks are not implemented properly, then it is possible for a
597 thread to overwrite others. We modified the algorithm [28] to capture this bug,
598 and modeled it as MPDA. BHIM found the bug with a witness of hole-bound 2.

599 • **Maze Example**. Finally we consider a robot navigating a maze, picking items;
600 an extended (from single to multiple stack) version of the example from [17]. In
601 the untimed setting, a witness for non-emptiness was obtained with hole-bound
602 0, while in the extension with time, the witness had a hole-bound 2, since the
603 satisfaction of time constraints required a longer witness.

604 **Results and Discussion**. The performance of BHIM is presented in Table 1 for
605 untimed examples and in Table 2 for timed examples. Apart from the results
606 in the tables, to check the robustness of BHIM wrt parameters like the number
607 of locations, transitions, stacks, holes and clocks (for TMPDA), we looked at
608 examples with an empty language, by making accepting states non-accepting in

Name	Locations	Transitions	Stacks	Holes	Time Empty (mili sec)	Time Witness (mili sec)	Memory(KB)
Bluetooth	57	96	2	0	157.9	7.1	7424
MultiProdCons	11	18	2	2	11.1	0.1	1796
dm-target	13	27	2	3	42.0	5.8	4476
Binary Search Tree	29	78	2	2	60.8	5.1	5143
untimed- L^{crit}	6	9	2	2	14.9	0.7	4692
untimed-Maze	9	12	2	0	12.0	0.2	3858
L^{bh} (from Sec. 2.1)	7	13	2	2	22.2	0.6	4404

Table 1. Experimental results: Time Empty and Time Witness column represents no. of milliseconds needed for emptiness checking and to generate witness respectively.

Name	Locations	Transitions	Stacks	Clocks	c_{max}	Aged(Y/N)	Holes	Time Empty(mili sec)	Time Witness (mili sec)	Memory(KB)
Bluetooth	57	96	2	0	2	Y	0	169.9	101.3	5248
L^{crit}	6	9	2	1	8	Y	2	1965.2	0.1	103396
L^{crit}	6	9	2	4	4	N	2	701664.2	0.1	1389432
Maze	9	12	2	2	5	Y	2	956.8	9.7	14554

Table 2. Experimental results of timed examples. The column c_{max} is defined as the maximum constant in the automaton, and Aged denotes if the stack is timed or not

the examples considered so far. This forces **BHIM** to explore all possible paths in the BFS tree, generating the lists at all nodes. The scalability of **BHIM** wrt all these parameters are in [11].

BHIM Vs. State of the art. What makes **BHIM** stand apart wrt the existing state of the art tools is that (i) none of the existing tools handle under approximations captured by bounded holes, (ii) none of the existing tools work with multiple stacks in the timed setting (even closed guards!). The state of the art research in underapproximations wrt untimed multistack pushdown systems has produced some amazing tools like GetaFix which handles multi-threaded programs with bounded context switches. While we have adapted some of the examples from GetaFix, the latest available version of GetaFix has some issues in handling those examples³. Likewise, SPADE, MAGIC and the counter implementation [27] are currently not maintained. This has come in the way of a performance comparison between **BHIM** and these tools. Indeed, most examples handled by **BHIM** correspond to non-context bounded, or non scope bounded, or timed languages which are beyond Getafix. For instance, the 2-hole bounded witness found by **BHIM** for the language $L_{20,10}$ ($m = 20, p = 10$) for the multi producer consumer case cannot be found by GetaFix/MAGIC/SPADE with less than 41 context switches. In the timed setting, the Maze example (TMPDA with 2 clocks, 2 timed stacks) has a 2 hole-bounded witness where the robot visits certain locations an equal number of times. The tool [17] cannot handle this example since it handles only one stack. Lastly, [17] cannot solve binary reachability with an empty stack unlike **BHIM**.

BHIM v2. The next version of **BHIM** will go symbolic, inspired from GetaFix. The current avatar of **BHIM** showcases the efficiency of fix-point techniques extended to larger bounded underapproximations; indeed going symbolic will make **BHIM** much more robust and scalable.

³ we did get in touch with the authors, who confirmed this

636 *Acknowledgements.* We would like to thank Gennaro Parlato for the discussions
637 we had on Getafix and for providing us benchmarks.

638 References

- 639 1. Salvatore La Torre, Parthasarathy Madhusudan, and Gennaro Parlato. A robust
640 class of context-sensitive languages. In *Logic in Computer Science, 2007. LICS*
641 *2007. 22nd Annual IEEE Symposium on*, pages 161–170. IEEE, 2007.
- 642 2. Salvatore La Torre, Parthasarathy Madhusudan, and Gennaro Parlato. The language
643 theory of bounded context-switching. In *Latin American Symposium on Theoretical*
644 *Informatics*, pages 96–107. Springer, 2010.
- 645 3. Salvatore La Torre, Margherita Napoli, and Gennaro Parlato. Scope-bounded
646 pushdown languages. *International Journal of Foundations of Computer Science*,
647 27(02):215–233, 2016.
- 648 4. Aiswarya Cyriac, Paul Gastin, and K Narayan Kumar. MSO decidability of multi-
649 pushdown systems via split-width. In *International Conference on Concurrency*
650 *Theory*, pages 547–561. Springer, Berlin, Heidelberg, 2012.
- 651 5. Salvatore La Torre and Margherita Napoli. Reachability of multistack pushdown
652 systems with scope-bounded matching relations. In *International Conference on*
653 *Concurrency Theory*, page 203–218. Springer, 2011.
- 654 6. Salvatore La Torre and Parlato Gennaro. Scope-bounded multistack pushdown
655 systems: Fixed-point, sequentialization, and tree-width. 2012.
- 656 7. Salvatore La Torre, Madhusudan Parthasarathy, and Gennaro Parlato. Analyzing
657 recursive programs using a fixed-point calculus. *ACM Sigplan Notices*, 44(6):211–
658 222, 2009.
- 659 8. Gaël Patin, Mihaela Sighireanu, and Tayssir Touili. Spade: Verification of
660 multithreaded dynamic and recursive programs. In *International Conference on*
661 *Computer Aided Verification*, pages 254–257. Springer, 2007.
- 662 9. Shaz Qadeer. The case for context-bounded verification of concurrent programs. In
663 *Model Checking Software, 15th International SPIN Workshop, Los Angeles, CA,*
664 *USA, August 10-12, 2008, Proceedings*, pages 3–6, 2008.
- 665 10. Sagar Chaki, Edmund Clarke, Nicholas Kidd, Thomas Reps, and Tayssir Touili.
666 Verifying concurrent message-passing C programs with recursive calls. In
667 *International Conference on Tools and Algorithms for the Construction and Analysis*
668 *of Systems*, page 334–349. Springer, 2006.
- 669 11. Akshay S, Gastin Paul, S Krishna, and Roychowdhury Sparsa. Supplementary
670 material: Revisiting under-approximate reachability in MPDA. Available at <https://cse.iitb.ac.in/~sparsa/bhim/>, 2019.
- 671 12. Mohamed Faouzi Atig. Model-checking of ordered multi-pushdown automata. *arXiv*
672 *preprint arXiv:1209.1916*, 2012.
- 673 13. Ahmed Bouajjani, Rachid Echahed, and Riadh Robbana. On the automatic
674 verification of systems with continuous variables and unbounded discrete data
675 structures. In *International Hybrid Systems Workshop*, pages 64–85. Springer, 1994.
- 676 14. Parosh Aziz Abdulla, Mohamed Faouzi Atig, and Jari Stenman. Dense-timed
677 pushdown automata. In *Proceedings of the 27th Annual IEEE Symposium on Logic*
678 *in Computer Science, LICS 2012, Dubrovnik, Croatia, June 25-28, 2012*, page
679 35–44, 2012.
- 680 15. Parosh Aziz Abdulla, Mohamed Faouzi Atig, and Jari Stenman. The minimal
681 cost reachability problem in priced timed pushdown systems. In *Language and*
682

- 683 *Automata Theory and Applications - 6th International Conference, LATA 2012, A*
684 *Coruña, Spain, March 5-9, 2012. Proceedings*, pages 58–69, 2012.
- 685 16. S. Akshay, Paul Gastin, and Shankara Narayanan Krishna. Analyzing Timed
686 Systems Using Tree Automata. *Logical Methods in Computer Science*, Volume 14,
687 Issue 2, May 2018.
- 688 17. S. Akshay, Paul Gastin, Shankara Narayanan Krishna, and Ilias Sarkar. Towards
689 an efficient tree automata based technique for timed systems. In *28th International*
690 *Conference on Concurrency Theory, CONCUR 2017, September 5-8, 2017, Berlin,*
691 *Germany*, pages 39:1–39:15, 2017.
- 692 18. Shaz Qadeer and Dinghao Wu. Kiss: keep it simple and sequential. *Acm sigplan*
693 *notices*, 39(6):14–24, 2004.
- 694 19. P Madhusudan and Gennaro Parlato. The tree width of auxiliary storage. In *ACM*
695 *SIGPLAN Notices*, volume 46, pages 283–294. ACM, 2011.
- 696 20. S. Akshay, Paul Gastin, Vincent Jugé, and Shankara Narayanan Krishna. Timed
697 systems through the lens of logic. In *34th Annual ACM/IEEE Symposium on Logic*
698 *in Computer Science, LICS 2019, Vancouver, BC, Canada, June 24-27, 2019*, pages
699 1–13, 2019.
- 700 21. Aiswarya Cyriac. *Verification of communicating recursive programs via split-width.*
701 *(Vérification de programmes récurrents et communicants via split-width)*. PhD thesis,
702 École normale supérieure de Cachan, France, 2014.
- 703 22. Rajeev Alur and Parthasarathy Madhusudan. Visibly pushdown languages. In
704 *Proceedings of the thirty-sixth annual ACM symposium on Theory of computing*,
705 pages 202–211. ACM, 2004.
- 706 23. Zhe Dang, Oscar H Ibarra, Tevfik Bultan, Richard A Kemmerer, and Jianwen Su.
707 Binary reachability analysis of discrete pushdown timed automata. In *International*
708 *Conference on Computer Aided Verification*, page 69–84. Springer, 2000.
- 709 24. Thomas H Cormen, Charles E Leiserson, Ronald L Rivest, and Clifford Stein.
710 *Introduction to algorithms*. MIT press, 2009.
- 711 25. Abraham Silberschatz, Greg Gagne, and Peter B Galvin. *Operating system concepts*.
712 Wiley, 2018.
- 713 26. Devendra Bhawe, Vrunda Dave, Shankara Narayanan Krishna, Ramchandra
714 Phawade, and Ashutosh Trivedi. A perfect class of context-sensitive timed languages.
715 In *International Conference on Developments in Language Theory*, pages 38–50.
716 Springer, Berlin, Heidelberg, 2016.
- 717 27. Matthew Hague and Anthony Widjaja Lin. Synchronisation- and reversal-bounded
718 analysis of multithreaded programs with counters. In *Computer Aided Verification*
719 *- 24th International Conference, CAV 2012, Berkeley, CA, USA, July 7-13, 2012*
720 *Proceedings*, page 260–276, 2012.
- 721 28. HT Kung and Philip L Lehman. Concurrent manipulation of binary search trees.
722 *ACM Transactions on Database Systems (TODS)*, 5(3):354–382, 1980.

Appendix

A Details for Section 2

A.1 Proposition 1

We use the notion of Tree Terms (TTs) [17] to compute the tree-width of a given graph. Where a minimal finite set of colors are used to color the vertices and then partition the graph in two partitions such that the cut vertices are colored. The aim of this approach is to decompose a graph to “atomic” tree terms. We cannot use a color more than once in a partition of graph, unless we *forget* it. This can be modeled as a game between two player, Adam and Eve. Where, Eve’s goal is to reach atomic terms with minimum finite number of colors, and Adam’s goal is to make Eve’s life difficult by choosing a more demanding partition.

To prove that the a model has bounded tree-width we will try to capture the runs of the model in terms of graphs (Multiply nested words [6]) and play the game mentioned above.

Tree Width of Hole Bounded Multistack Pushdown Automaton

We will capture the behaviour (any run ρ) of K -hole bounded multistack pushdown systems as a graph G where, every node represents a transition $t \in \Delta$ and the edge between the nodes can be of the following types.

- Linear order \preceq between the transitions gives the order in which the transitions are fired in the system. We will use \preceq^+ to represent transitive closure of \preceq .
- The other type of edges represent the push pop relation between two transitions. Which means, if a transition t_1 have a push operation in the stack i and transition t_2 has the corresponding pop of the stack i , matching the push on stack i of transition t_1 , then we have an edge $t_1 \curvearrowright^s t_2$ between them, which will represent the push-pop relation.

To prove that the tree width of the class of graph G is bounded, we will use coloring game [17] and show that we need bounded number of colors to split any graph $g \in G$ to atomic tree terms.

Eve will start from the right most node of the graph by coloring it. The last node of the graph can be any one of the following,

- End point of a well-nested sequence
- Pop transition t_{pp} of stack i , such that, the push t_{ps} is coming from nearest *hole* of stack i .

1. If the endpoint colored is the end point of a well-nested sequence then Eve can remove the well-nested sequence by adding another color to the first point of the well-nested sequence.

If we look at the well-nested part, using just one more color we can split it to atomic tree terms [17].

But the other part still remains a graph of class G so Adam will choose this partition for Eve to continue the coloring game.

762 2. If the last point of the graph G is a pop point t_{pp} as discussed earlier, then
763 the corresponding push t_{ps} can come from an open *hole* or a closed *hole*.
764 – If it is coming from a closed *hole* then, Eve will add color to the
765 corresponding push t_{ps} along with the transition t_q such that, $t_{ps} \preceq^+ t_q$
766 and $t_{ps} - t_q$ is a well-nested sequence, which forms an atomic hole segment
767 $(\uparrow ws)$ where, \uparrow represents the push pop edge $t_{ps} \curvearrowright^s t_{pp}$ and ws represents
768 the well-nested sequence $t_{ps} - t_q$. This operation requires 2 colors. Please
769 note that, the right end of the *hole* which got colored after removal of
770 $t_{ps} - t_q$ is another push of the *hole*, because *hole* are defined as a sequence
771 $(\uparrow ws)^+$.
772 – If the push is coming from open *hole* then the push transition t_{ps} is already
773 colored from previous operation as discussed above, hence Eve will add
774 another color $t_{q'}$ to mark the next well-nested sequence $t_{ps} - t_{q'}(ws')$ in
775 the right of t_{ps} . Now, Eve can remove the stack edge $t_{pp} \curvearrowright t_{ps}$ along
776 with the well-nested sequence ws' . This operation widens the *hole*.
777 In both the above operations, the graph has two components one with a stack
778 edge $t_{pp} \curvearrowright t_{ps}$ and another one with a well-nested sequence. Which require
779 at most 1 color extra to split into atomic tree terms. On the remaining part
780 Eve will continue playing the game from right most point.

781 Here, we claim that at any point of time of the coloring game, there will be
782 $2K + 2$ active colors for $K \geq 1$ and $K \in \mathbb{N}$. Every step of the game splits the
783 graph in two part, and one part always can be split into atomic tree terms with
784 at most 3 colors. The remaining part will require at most 2 colors for every open
785 *hole* in the left of the right most point of the graph. As the number of open *hole*
786 is bounded by K , so we can not have more than K open *holes* in the left of any
787 point. So, $2K$ colors to mark the *holes*. So, total number of colors needed to
788 break any such graph to atomic tree terms is $2K + 4$.

789 A.2 Proposition 2

790 We describe the missing details in proposition 2.

- 791 1. L^{bh} cannot be accepted by any K -bounded phase MPDA.
792 Recall that, $L^{bh} = \{a^n b^n (a^{q_i} c^{q_i+1} b^{q'_j} d^{q'_j+1})^n \mid n, q_i, q'_j \in \mathbb{N} \forall i, j \in [n]\}$, and a, b
793 represents push in stack 1,2 respectively and c, d represents the corresponding
794 pops from stack 1,2. For all m , consider the word $w_1 = a^m b^m (a^l c^{l+1} b^{l'} d^{l'+1})^m$.
795 Here, clearly the number of phases is $K = 2m$. Now if w_1 is accepted by some
796 phase bounded MPDA M then it must have $2m$ as the bound on the phases
797 which will not be sufficient to accept $w_2(a^{m+1} b^{m+1} (a^l c^{l+1} b^{l'} d^{l'+1})^{m+1}) \in$
798 L^{bh} .
799 2. $L' = \{(ab)^n c^n d^n \mid n \in \mathbb{N}\}$ cannot be accepted by any K -hole bounded MPDA.
800 For any $m \in \mathbb{N}$ assume a word $w_1 = (ab)^m c^m d^m \in L'$, where a, b represents
801 push in stack 1,2 respectively and c, d represents the corresponding pops
802 from stack 1,2. Clearly, this can be accepted by a bounded *hole* multistack
803 pushdown automata M with bound $= 2m$. Now if L' is accepted by M then

804 it must also accept, $w_2 = (ab)^{m+1}c^{m+1}d^{m+1}$. However, the number of *holes*
805 required to accept w_2 is $2(m+1) > 2m$. This contradicts the assumption
806 that M accepts the language.

807 B Details for Section 3

808 In this section, we provide all the subroutines mentioned in Section 3 and used
809 in Algorithm 1 for MPDA.

810 We start by presenting Algorithm 2 which computes the well-nested reachability
811 relation, i.e., it computes the set WR of all pairs of states (s, s') such that there
is a well-nested sequence from s to s' . The proof of correctness of this algorithm

Algorithm 2: Well Nested Reachability

```

1 Function WellNestedReach( $M = (S, \Delta, s_0, S_f, n, \Sigma, \Gamma)$ ):
   Result:  $WR := \{(s, s') \mid s' \text{ is reachable from } s \text{ via a well-nested sequence}\}$ 
2    $\mathcal{R}_c := \{(s, s) \mid s \in S\};$ 
3   forall  $(s_1, \text{op}, a, s_2) \in \Delta$  with  $\text{op} = \text{nop}$  do
4      $\mathcal{R}_c := \mathcal{R}_c \cup \{(s_1, s_2)\};$   $\backslash \backslash$  Transitions with nop operation
5    $\mathcal{R}_c := \text{TransitiveClosure}(\mathcal{R}_c);$   $\backslash \backslash$  Using Floyd-Warshall Algorithm
6   while True do
7      $WR := \mathcal{R}_c;$ 
8     forall  $(s, \downarrow_i(\alpha), a, s_1) \in \Delta$  do
9       forall  $(s_1, s_2) \in WR$  do
10        forall  $(s_2, \uparrow_i(\alpha), a, s') \in \Delta$  do
11           $\mathcal{R}_c := \mathcal{R}_c \cup \{(s, s')\};$   $\backslash \backslash$  Wrap well-nested sequence with
             matching push-pop
12         $\mathcal{R}_c := \text{TransitiveClosure}(\mathcal{R}_c);$ 
13        if  $\mathcal{R}_c \setminus WR = \emptyset$  then
14          break;  $\backslash \backslash$  Break when no new well-nested sequence added
15 return  $WR;$ 

```

812 (and thus Lemma 1) is easy to see. First, line 5 the set \mathcal{R}_c contains all pairs (s, s')
813 such that s' is reachable from s in the MPDA without using the stack. Then
814 for every push transition from a state s we check in lines 8-11 whether there
815 is an (already computed) well-nested sequence that can reach a state s' with
816 a corresponding pop transition and if so we add (s, s') . We take the transitive
817 closure and repeat this process, hence guaranteeing that at fixed point we will
818 have all well-nested pairs, i.e., WR .
819

820 *Details of Algorithm 3* For a given list μ Algorithm 3 tries to extend the list μ
821 by adding a hole of a stack i . This is achieved by checking the last state s_{last} the
822 list μ and finding all possible hole in HS_i that start with s_{last} and appending
823 the hole followed by a suitable well-nested sequence to μ .

Algorithm 3: AddHole

```
1 Function AddHolei( $\mu$ ,  $HS_i$ ):  
   Result: Set, a set of lists.  
2    $Set := \emptyset$ ;  
3   forall  $(i, s, s') \in HS_i$  with  $s = \text{last}(\mu)$  do  
4      $\mu' := \text{copy}(\mu)$ ;  $\backslash \backslash$  Create a copy of the list  $\mu$   
5      $\text{trunc}(\mu')$ ;  $\backslash \backslash$  trunc( $\mu$ ) is defined as  $\text{remove}(\text{last}(\mu))$   
6      $\mu'.\text{append}[(i, s, s'), s']$ ;  $\backslash \backslash$  Append to the list  $\mu'$   
7      $\mu'.\text{NumberOfHoles} := \mu.\text{NumberOfHoles} + 1$ ;  
8      $Set := Set \cup \{\mu'\}$ ;  
9 return  $Set$ ;
```

Algorithm 4: Extend with a pop

```
1 Function AddPopi( $\mu$ ,  $M = (\mathcal{S}, \Delta, s_0, \mathcal{S}_f, n, \Sigma, \Gamma)$ ,  $AHS_i$ ,  $HS_i, WR$ ):  
   Result: Set, a set of lists  
2    $Set := \emptyset$ ;  
3    $(i, s_1, s_3) := \text{lastHole}_i(\mu)$ ;  $\backslash \backslash$  Get the last open hole of stack  $i$   
4   forall  $(i, s_1, s_2) \in HS_i$ ,  $(s_2, \alpha, s_3) \in AHS_i$ ,  $(s, \uparrow_i(\alpha), s') \in \Delta$ ,  $s = \text{last}(\mu)$   
   and  $(s', s'') \in WR$  do  
5      $\mu' := \text{copy}(\mu)$ ;  
6      $\text{trunc}(\mu')$ ;  
7      $\mu'.\text{append}(s'')$ ;  
8     if  $(s_1 = s_2)$  then  
9        $\mu'' := \text{copy}(\mu)$ ;  
10       $\text{trunc}(\mu'')$ ;  
11       $\mu''.\text{append}(s'')$ ;  
12       $\mu''.\text{remove}((i, s_1, s_3))$ ;  $\backslash \backslash$  Remove the hole  $(i, s_1, s_2)$  from the  
       list  $\mu''$   
13       $\mu''.\text{NumberOfHoles} := \mu.\text{NumberOfHoles} - 1$ ;  
14       $Set := Set \cup \{\mu''\}$ ;  
15       $\mu'.\text{replace}((i, s_1, s_3), \text{by } (i, s_1, s_2))$ ;  $\backslash \backslash$  Replace bigger hole  
        $(i, s_1, s_3)$  by new smaller hole  $(i, s_1, s_2)$   
16       $Set := Set \cup \{\mu'\}$ ;  
17 return  $Set$ ;
```

824 *Details of Algorithm 4* For a given list μ this algorithm tries to extend μ with a
825 pop operation. The algorithm starts with extracting the last hole(H_i) of stack
826 i . Due to the well-nested property, the pop (which is not part of a well-nested
827 sequence) must be matched with the first pending push in the last hole of stack
828 i in μ . Then the algorithm checks for all atomic hole-segments AHS_i and hole-
829 segments HS_i s of the stack i , such that, the hole H_i can be partitioned in HS_i
830 and AHS_i . Then the push in AHS_i is matched with the matched pop operation
831 and the hole is now shrunk into HS_i . So, the algorithm replaces H_i with HS_i . If
832 the H_i is same as some AHS_i then, the hole can be closed and hence it removes
833 the hole from the list. In this case it also reduces the count of the number of

834 holes in the list. Note that without the pre-computation of AHS_i and HS_i this
 835 part of the algorithm is fairly difficult. Using the pre-computation allow us to use
 836 simple table look ups when the states are known, this takes only constant time.

837 C Details for Section 4

838 The algorithm for witness generation, as discussed in the main part of the paper,
 839 does a backtracking on the BFS tree. When we encounter a node in the BFS
 840 tree extending the list with a pop, creating a hole, we use the last state in
 841 the list, the transition information from the node, and the witness stack for
 842 backtracking. During the backtracking we also need to know the sequence of
 843 transitions responsible for the well-nested sequences, which can be generated
 844 using the Algorithm 5. The backtracking Algorithm 6 is discussed in the following
 845 example.

Algorithm 5: Well-nested witness generation for MPDA

```

1 Function  $\text{Witness}(s_1, s_2, M = (\mathcal{S}, \Delta, s_0, \mathcal{S}_f, n, \Sigma, \Gamma), WR)$ :
   Result: A sequence of transitions for a run resulting the well-nested
           sequence  $WR$ 
2   if  $s_1 == s_2$  then
3     return  $\epsilon$ ;
4   if  $\exists t = (s_1, \text{nop}, a, s_2) \in \Delta$  then
5     return  $t$ ;
6   forall  $s', s'' \in \mathcal{S}$  do
7     if  $((s_1 \neq s') \vee (s'' \neq s_2)) \wedge (s', s'') \in WR \wedge \exists t = (s_1, \downarrow_i(\alpha), a, s') \in \Delta \wedge$ 
        $\exists t_2 = (s'', \uparrow_i(\alpha), a', s_2) \in \Delta$  then
8        $\text{path} = \text{Witness}(s', s'', M, WR)$ ;
9       return  $t.\text{path}.t_2$ ;
10  forall  $s \in \mathcal{S}$  do
11    if  $(s \neq s_1 \vee s \neq s_2) \wedge (s, s_1) \in WR \wedge (s, s_2) \in WR$  then
12       $\text{path1} = \text{Witness}(s_1, s, M, WR)$ ;
13       $\text{path2} = \text{Witness}(s, s_2, M, WR)$ ;
14    return  $\text{path1}.\text{path2}$ ;

```

846 An Illustrating Example for Witness Generation

847 We illustrate the multistack case on an example. Note that in figures illustrating
 848 examples, we use colored uparrows and downarrows with subscript for stacks,
 849 and a superscript i representing the i th push or pop of the relevant colored stack.

850 Assume that the path we obtain on back tracking is the reverse of Figure 3.
 851 Holes arising from pending pushes of stack 1 are red holes, and those from stack
 852 2 are blue holes in the figure. We have two red holes: the first red hole has a left
 853 end point \downarrow_1^1 , and right end point ws_3 . The second red hole has a left end point

Algorithm 6: Non-well-nested witness generation for MPDA

```

1 Function HoleWitness( $\mu, M = (\mathcal{S}, \Delta, s_0, \mathcal{S}_f, n, \Sigma, \Gamma), WR, AHS_i, HS_i$ ):
   Result: A sequence of transitions for an accepting run
2 global WitnessStacks =  $\{St_i \mid i \in [n]\}$ ; Witness stacks for every
   stack i
3  $\mu_p = \text{Parent}(\mu)$ ; Parent function returns the parent node of  $\mu$  in
   the BFS exploration tree
4  $op_\mu = \text{ParentOp}(\mu)$ ; ParentOp function returns the operation
   that extends  $\text{Parent}(\mu)$  to  $\mu$  in the BFS exploration tree
5 if  $op_\mu == \text{ExtendByPop}_i(\uparrow_i \alpha.wr_{pop}) \wedge wr_{pop} \in WR$  then
6    $(i, s_1, s_2) = \text{lastHole}_i(\mu_p)$ ;
7   if  $(s_i, \alpha, s_2) \in AHS_i \wedge (s_1, \alpha, s_2) = \downarrow_i(\alpha).wr_{push} \wedge wr_{push} \in WR$  then
8      $push(St_i, \#)$ ;
9      $list = \text{Witness}(wr_{push})$ ;
10     $\forall t \in list, push(St_i, t)$ ;
11     $push(St_i, \downarrow_i(\alpha))$ ;
12     $list_{pop} = \text{Witness}(wr_{pop})$ ;
13    return HoleWitness( $\mu_p$ ). $\uparrow_i(\alpha).list_{pop}$ ;
14  else if
     $(s_i, \alpha, s_2) \notin AHS_i \wedge (i, s_i, s_2) = (s_i, \alpha, s_3).(i, s_3, s_2) \wedge (s_1, \alpha, s_3) \in$ 
     $AHS_i \wedge (i, s_3, s_2) \in HS_i \wedge (s_1, \alpha, s_3) = \downarrow_i(\alpha).wr_{push} \wedge wr_{push} \in WR$ 
    then
15     $list = \text{Witness}(wr_{push})$ ;
16     $\forall t \in list, push(St_i, t)$ ;
17     $push(St_i, \downarrow_i(\alpha))$ ;
18     $list_{pop} = \text{Witness}(wr_{pop})$ ;
19    return HoleWitness( $\mu_p$ ). $\uparrow_i(\alpha).list_{pop}$ ;
20  if  $op_\mu == \text{ExtendByHole}_i$  then
21     $list = \epsilon$ ;
22    while  $pop(St_i) \neq \#$  do
23       $list = list.pop(St_i)$ ;
24    return HoleWitness( $\mu_p$ ). $list$ ;

```

854 \downarrow_1^4 , and right end point \downarrow_1^5 . The blue hole has left end point \downarrow_2^1 and right end
 855 point ws_4 .

- 856 1. From the final configuration s_f , on backtracking, we obtain the pop operation
 857 (\uparrow_1^1) . By the fixed-point algorithm, this operation closes the first red hole,
 858 matching the first pending push \downarrow_1^1 . In the BFS exploration tree, the parent
 859 node has the red atomic hole consisting of just the \downarrow_1^1 . Notice also that,
 860 in the parent node, this is the only red hole, since the second red hole in
 861 Figure 3 is closed, and hence does not exist in the parent node. We use two
 862 witness stacks, a red witness stack and a blue witness stack to track the
 863 information with respect to the red and blue holes. On encountering a pop
 864 transition closing a red hole, we populate the red witness stack with (i) a
 865 barrier signifying closure of a red hole, and (ii) the matching push transition
 866 \downarrow_1^1 .

9. Further backtracking leads us to the parent obtained by extending with the well-nested sequence ws_1 . We retrieve the transitions in ws_1 using Algorithm 5. The last backtracking lands us at the root $[s_0]$ and we are done.

D Details for Section 5

This part of the appendix is devoted to extending our algorithms for reachability and witness generation. We start by defining timed multistack push down automata. Then, Appendix E details the (binary) reachability and algorithms therein, whereas Appendix F describes the generation of a witness for TMPDA.

Timed Multi-stack Pushdown Automata (TMPDA)

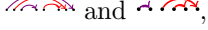
For $N \in \mathbb{N}$, we denote the set of numbers $\{1, 2, 3 \dots N\}$ as $[N]$. \mathcal{I} denotes the set of closed intervals $\{I \mid I \subseteq \mathbb{R}_+\}$, such that the end points of the intervals belong to \mathbb{N} . \mathcal{I} also contains a special interval $[0, 0]$. We start by defining the model of timed multi-pushdown automata.

Definition 3. A *Timed Multi-pushdown automaton* (TMPDA [16]) is a tuple $M = (\mathcal{S}, \Delta, s_0, \mathcal{S}_f, \mathcal{X}, n, \Sigma, \Gamma)$ where, \mathcal{S} is a finite non-empty set of locations, Δ is a finite set of transitions, $s_0 \in \mathcal{S}$ is the initial location, $\mathcal{S}_f \subseteq \mathcal{S}$ is a set of final locations, \mathcal{X} is a finite set of real valued variables known as clocks, n is the number of (timed) stacks, Σ is a finite input alphabet, and Γ is a finite stack alphabet which contains \perp . A transition $t \in \Delta$ can be represented as a tuple $(s, \varphi, \text{op}, a, R, s')$, where, $s, s' \in \mathcal{S}$ are respectively, the source and destination locations of the transition t , φ is a finite conjunction of closed guards of the form $x \in I$ represented as, $(x \in I' \wedge y \in I'' \dots)$ for $x, y \in \mathcal{X}$ and $I', I'' \in \mathcal{I}$, $R \subseteq \mathcal{X}$ is the set of clocks that are reset, $a \in \Sigma$ is the label of the transition, and op is one of the following stack operations (1) **nop**, or no stack operation, (2) $(\downarrow_i \alpha)$ which pushes $\alpha \in \Gamma$ onto stack $i \in [n]$, (3) $(\uparrow_i^I \alpha)$ which pops stack i if the top of stack i is $\alpha \in \Gamma$ and the time elapsed from the push is in the interval $I \in \mathcal{I}$.

A *configuration* of TMPDA is a tuple $(s, \nu, \lambda_1, \lambda_2, \dots, \lambda_n)$ such that, $s \in \mathcal{S}$ is the current location, $\nu: \mathcal{X} \rightarrow \mathbb{R}$ is the current clock valuation and $\lambda_i \in (\Gamma \times \mathbb{R})^*$ represents the current content of i^{th} stack as well as the *age* of each symbol, i.e., the time elapsed since it was pushed on the stack. A pair (s, ν) , where s is a location and ν is a clock valuation is called a *state*.

The semantics of the TMPDA is defined as follows: a run σ is a sequence of alternating time elapse and discrete transitions from one configuration to another. The time elapses are non-negative real numbers, and, on discrete transitions, the valuation ν of the current configuration is checked to see if the clock constraints are satisfied; likewise, on a pop transition, the age of the symbol popped is checked. Projecting out the operations of a single stack from σ results in a well-nested sequence. A run is accepting if it starts from the initial state with all clocks set to 0, and reaches a final state with all stacks empty. The language accepted by a TMPDA is defined as the set of timed words generated by the accepting runs

945 of the TMPDA. Since the reachability problem for TMPDA is Turing complete
 946 (this is the case even without time), we consider under-approximate reachability.

947 A sequence of transitions is said to be **complete** if each push has a matching
 948 pop and vice versa. A sequence of transitions is said to be **well-nested**, denoted
 949 ws , if it is a sequence of **nop**-transitions, or a concatenation of well-nested
 950 sequences ws_1ws_2 , or a well-nested sequence surrounded by a matching push-pop
 951 pair $(\downarrow_i\alpha)ws(\uparrow_i^I\alpha)$. If we visualize this by drawing edges between pushes and
 952 their corresponding pops, well-nested sequences have no crossing edges, as in
 953 , where we have two stacks, depicted with red and violet edges.
 954 We emphasize that a well-nested sequence can have well-nested edges from any
 955 stack. In a sequence σ , a push (pop) is called a **pending** push (pop) if its
 956 matching pop (push) is not in the same sequence σ . For TMPDA every sequence
 957 also carries total time elapsed during the sequence, this is helpful to check stack
 958 constraints, and it is sufficient to store time till the maximum stack constraint,
 959 i.e., the maximum constant value that appeared in the stack constraints.

960 Tree Width of Bounded Hole TMPDA

961 We will capture the behaviour (any run ρ) of K -hole bounded multistack
 962 pushdown systems as a graph G where, every node represents a transition
 963 $t \in \Delta$ and the edge between the nodes can be of three types.

- 964 – Linear order (\preceq) between the transition which gives the order in which the
 965 transitions are fired. We will use \preceq^+ to represent transitive closure of \preceq .
- 966 – Timing relations $\curvearrowright^{c \in I} \in \preceq^+ \forall c \in \mathcal{X}$ and $I \in \mathcal{I}$ such that, $t_1 \curvearrowright^{c \in I} t_2$ if and
 967 only if the clock constraint $c \in I$ is checked in the transition t_2 and $t_1 \preceq^+ t_2$
 968 has the latest reset of clock c with respect to t_2 .
- 969 – The other type of edges represent the push pop relation between two
 970 transitions. Which means, if a transition t_1 have a push operation in any one
 971 of the stack i and transition t_2 has pop transition of the stack i which matches
 972 with the push transition at t_1 , then we have an edge $t_1 \curvearrowright^s t_2$ between them,
 973 which will represent the stack edge.

974 To prove that the tree width of the class of graph G is bounded, we will use
 975 coloring game [17] and show that we need bounded number of colors to split any
 976 graph $g \in G$ to atomic tree terms.

977 Eve will start from the right most node of the graph by coloring it. The last
 978 node of the graph can be any one of the following,

- 979 – End point of a well-nested sequence
- 980 – Pop transition t_{pp} of stack i , such that, the push t_{ps} is coming from nearest
 981 hole of stack i .

- 982 1. If the end point colored is the end point of a well-nested sequence then Eve
 983 can remove the well-nested sequence by adding another color to the first
 984 point of the well-nested sequence. But, there may be some transitions t in the
 985 well-nested sequence with clock constraints $c \in I$ such that, the recent reset

of the clock c , with respect to t is in the left of the well nested sequence. In order to remove the well-nested sequence she have to color the nodes which represent the transitions with recent reset points of the clocks $c \in \mathcal{X}$. This step require at most $|\mathcal{X}|$ colors. Now, she can split the graph in two parts, one of them will be well-nested with two end points colored. Also, the clock constraint edges, which are coming from the left of the well-nested sequence are hanging in the left, are colored. There can be at most $|\mathcal{X}|$ hanging colored points possible in the left of the well-nested sequence. The other part will be the remaining graph with the right most point colored along with the colored recent reset points on the left of right most colored point. which are also the hanging points of the previous partition.

If we look at the well-nested part with hanging clock edges, using just one more color we can split it to atomic tree terms [17].

But the other part still remains a graph of class G so Adam will choose this partition for Eve to continue the coloring game.

2. If the last point of the graph G is a pop point t_{pp} as discussed earlier, then the corresponding push t_{ps} can come from a open *hole* or a closed *hole*.
 - If it is coming from a closed *hole* then, Eve will add color to the corresponding push t_{ps} along with the transition t_q such that, $t_{ps} \preceq^+ t_q$ and $t_{ps} - t_q$ is a well-nested sequence, which forms a atomic hole segment $(\uparrow ws)$ where, \uparrow represents the push pop edge $t_{ps} \curvearrowright^s t_{pp}$ and ws represents the well-nested sequence $t_{ps} - t_q$. But just as we discussed in previous scenario of removing well-nested sequence, there may be some clock constraint $c \in \mathcal{X}$ in the well-nested sequence ws such that the transition with the recent resets are from the left of $(\uparrow ws)$ and without coloring them Eve can not remove the $(\uparrow ws)$. Similarly, there may be some clock resets inside $\uparrow ws$ from which there are clock constraint edges are going to the right of $\uparrow ws$. Eve has to color all those points inside the $\uparrow ws$ which corresponds to those clock reset points in $\uparrow ws$. So, she have to color at most $2|\mathcal{X}|$ reset points to remove the stack edge $t_1 \curvearrowright t_2$ along with the well-nested sequence $t_{ps} - t_q(\uparrow ws)$, which makes the closed *hole* open with colors in both ends of hole and at most $|\mathcal{X}|$ colors in the left of the hole and at most $|\mathcal{X}|$ colored hanging points inside the *hole*. This operation requires $2 + 2|\mathcal{X}|$ more colors. Please note that, the right end of the *hole* which got colored after removal of $t_{ps} - t_q$ is another push of the *hole*, because *hole* are defined as a sequence $(\uparrow ws)^+$.
 - If the push is coming from open *hole* then the push transition t_{ps} must be colored from previous operation as discussed above, hence Eve will add another color $t_{q'}$ to mark the next well-nested sequence $t_{ps} - t_{q'}(ws')$ in the right of t_{ps} . But, similar to above section here also there may be some clock resets of clock $i \in \mathcal{X}$ inside the ws' which is being checked in the right of the ws' . These reset points can be at most $|\mathcal{X}|$ and needs $|\mathcal{X}|$ colors. Now, Eve can remove the stack edge $t_{pp} \curvearrowright t_{ps}$ along with the well-nested sequence ws' . This operation widens the *hole*. Note that at any point of the game, hanging clock reset points inside the *hole* and in left side of hole is bounded by $|\mathcal{X}|$. This operation requires at most

1032 $1 + |\mathcal{X}|$ colors but subsequent application of this operation can reuse
 1033 colors.

1034 In both the above operations, we can split the graph in two parts, one with
 1035 a stack edge $t_{pp} \curvearrowright t_{ps}$ and a well-nested sequence, with at most $|\mathcal{X}|$ hanging
 1036 points for each clock in the left of the t_{pp} and at most $|\mathcal{X}|$ colors inside the
 1037 ws . which require at most 1 color to split into atomic tree terms without any
 1038 extra colors. On the remaining part Eve will continue playing the game from
 1039 right most point.

1040 Here, we claim that at any point of time of the coloring game, there will be
 1041 $2K + (2K + 1)|\mathcal{X}| + 2$ active colors for $K \geq 1$ and $K \in \mathbb{N}$. Every step of the
 1042 game splits the graph in two part, and one part always can be split into atomic
 1043 tree terms with at most $2|\mathcal{X}| + 3$ colors. The remaining part will require $2 + 2|\mathcal{X}|$
 1044 colors for every open *hole* in the left of the right most point of the graph. As the
 1045 number of open *hole* is bounded by K , so we can not have more than K open
 1046 *holes* in the left of any point. So, $2K + 2K|\mathcal{X}|$ colors to mark the *holes*, $1 + |\mathcal{X}|$
 1047 for the right most point and recent reset points with respect to the right most
 1048 point, $1 + |\mathcal{X}|$ for coloring the well-nested sequence after a matched push and the
 1049 possible reset points inside the well-nested sequence, but we will need to color
 1050 such well-nested sequence once at any point of time, which gives a total color of
 1051 $2K(|\mathcal{X}| + 1) + 2(|\mathcal{X}| + 1) = (2K + 2)(|\mathcal{X}| + 1)$.

1052 E Reachability in TMPDA

1053 In this section, we discuss how the BFS tree exploration extends in the timed
 1054 setting. To begin, we talk about how a list at any node in the tree looks like.

1055 Representation of Lists for BFS Tree

1056 Each node of the BFS tree stores a list of bounded length. A list is a sequence of
 1057 states (s, ν) separated by time elapses (t) , representing a K -hole bounded run in
 1058 a concise form. The simplest kind of list is a single state (s, ν) or a well-nested
 1059 sequence (s, ν, t, s_i, ν_i) with time elapse t . Note that because of time constraints
 1060 we need to store total time elapsed to reach one state from another. This is
 1061 why we are keeping a time stamp between two states. Recall, the hole in MPDA
 1062 is defined as a tuple (i, s, s') . For TMPDA we need to store total time elapsed
 1063 in the hole as well, so it can be represented as a tuple $H = (i, s, \nu, s', \nu', t_h)$,
 1064 where, t_h is the time elapse in the hole and $(s, \nu), (s', \nu')$ being the end states
 1065 of the hole. Also, the maximum possible value of time stamp is bounded by the
 1066 maximum integer value in the constraints (both pop and clock). So, the total
 1067 possible values that the variable t_i can take is also bounded. Let H, t represent
 1068 respectively holes (of some stack) and time elapses. A list with holes has the form
 1069 $(s_0, \nu_0).t.(H)^*(H.t.(s', \nu'))$. For example, a list with 3 holes of stacks i, j, k is

1070 $[(s_0, \nu_0), t_1, (i, s_1, \nu_1, s_2, \nu_2, t_2), t_3, (j, s_3, \nu_3, s_4, \nu_4, t_4), t_5, (k, s_5, \nu_5, s_6, \nu_6, t_6), t_7, (s_7, \nu_7)]$

Algorithm 7: Algorithm for Emptiness Checking of hole bounded TMPDA

```

1 Function IsEmptyTimed( $M = (\mathcal{S}, \Delta, s_0, \mathcal{S}_f, \mathcal{X}, n, \Sigma, \Gamma), K$ ):
   Result: True or False
2    $WRT := \text{WellNestedReachTimed}(M)$ ;  $\backslash \backslash$  Solves binary reachability for pushdown system
3   if some  $(s_0, \nu_0, t, s_1, \nu_1) \in WRT$  with  $s_1 \in \mathcal{S}_f$  then
4     return False;
5   forall  $i \in [n]$  do
6      $AHST_i := \emptyset$ ;
7     forall  $(s, \phi, \downarrow_i(\alpha), \rho, a, s_1) \in \Delta, \nu \models \phi$ , and  $\nu_1 = \rho[\nu]$  do
8       forall  $(s_1, \nu_1, t, s', \nu') \in WRT$  do
9          $AHST_i := AHST_i \cup (i, s, \nu, \alpha, s', \nu', t)$ ;
10       $Set_i := \{(s, \nu, t, s', \nu') \mid \exists \alpha(i, s, \nu, \alpha, s', \nu', t) \in AHST_i\}$ ;
11       $HS_i := \{(i, s, \nu, s', \nu', t) \mid (s, \nu, t, s', \nu') \in \text{TransitiveClosure}(Set_i)\}$ ;
12       $\mu := [s_0, \nu_0]$ ;
13       $\mu.\text{NumberOfHoles} := 0$ ;
14       $SetOfLists_{new} := \{\mu\}, SetOfLists_{old} := \emptyset$ ;
15      while  $SetOfLists_{new} \setminus SetOfLists_{old} \neq \emptyset$  do
16         $SetOfLists_{diff} := SetOfLists_{new} \setminus SetOfLists_{old}$ ;
17         $SetOfLists_{old} := SetOfLists_{new}$ ;
18        forall  $\mu' \in SetOfLists_{diff}$  do
19          if  $\mu'.\text{NumberOfHoles} < K$  then
20            forall  $i \in [n]$  do
21               $SetOfLists_h := \text{AddHoleTimed}_i(\mu', HS_i)$ ;  $\backslash \backslash$  Add hole for stack i
22              forall  $\mu_2 \in SetOfLists_h$  do
23                 $SetOfLists_{new} := SetOfLists_{new} \cup \mu_2$ ;
24            if  $\mu'.\text{NumberOfHoles} > 0$  then
25              forall  $i \in [n]$  do
26                 $SetOfLists_p := \text{AddPopTimed}_i(\mu', M, AHST_i, HS_i, WRT)$ ;  $\backslash \backslash$  Add pop
27                for stack i
28                forall  $\mu_3 \in SetOfLists_p$  do
29                  if  $\mu_3.\text{last} \in \mathcal{S}_f$  and  $\mu_3.\text{NumberOfHoles} = 0$  then
30                    return False;  $\backslash \backslash$  If reached destination state
31                     $SetOfLists_{new} := SetOfLists_{new} \cup \mu_3$ ;
32      return True;

```

Algorithm 8: States

```

1 Function States( $M = (\mathcal{S}, \Delta, s_0, \mathcal{S}_f, \mathcal{X}, n, \Sigma, \Gamma)$ ):
   Result:  $F$ 
2    $F := \{(s, \nu) \mid \forall s \in \mathcal{S} \wedge \forall c \in \mathcal{X}, \nu[c] \leq \max(c) + 1\}$ ;
3   return  $F$ ;

```

1071 **Algorithms for TMPDA**

1072 The function TimeElapse returns the states which are reachable from the state
1073 (s_1, ν_1) via time elapse. It also stores the total time elapsed to reach the state.
1074 This function is only useful for timed systems.

Algorithm 9: Time Elapse

```
1 Function TimeElapse( $(s_1, \nu_1)$ ):  
   Result: Set  
2    $Set := \emptyset$ ;  
3    $t := 0$ ;  
4   while  $t \leq c_{max}$  do  
5      $\forall i \in X : \nu_2[i] := \text{Min}(\nu_1[i] + t, c_i)$ ;  
6      $Set := Set \cup (s_1, \nu_1, t, s_1, \nu_2)$  ;  
7      $t := t + 1$ ;  
8 return  $Set$ ;
```

Algorithm 10: Well Nested Reach Timed

```
1 Function WellNestedReachTimed( $M = (S, \Delta, s_0, S_f, \mathcal{X}, n, \Sigma, I)$ ):  
   Result:  $\text{WRT} := \{(s, \nu, t, s', \nu') \mid (s', \nu') \text{ is reachable from } (s, \nu) \text{ by time elapse } t \text{ via a well-nested sequence}\}$   
2    $F = \text{States}(M)$ ;  
3    $Set = \{(s, \nu, p, s, \nu) \mid (s, \nu) \in F\}$ ;  
4   forall  $(s, \nu) \in F$  do  
5      $Set = Set \cup \text{TimeElapse}((s, \nu))$ ;  
6     forall  $(s, \varphi, \text{nop}, a, R, s') \in \Delta$  with  $\nu \models \phi$  do  
7        $Set := Set \cup (s, \nu, 0, s', R[\nu])$   
8    $\mathcal{R}_{tc} = \text{TransitiveClosureTimed}(Set)$ ;  
9   while True do  
10     $\text{WRT} := \mathcal{R}_{tc}$ ;  
11    forall  $(s, \phi_1, \downarrow_i(\alpha), \rho_1, a, s_1) \in \Delta$  and  $(s, \nu) \in F$  with  $\nu \models \phi_1$  do  
12      forall  $(s_1, \rho_1[\nu], t, s_2, \nu_2) \in \mathcal{R}_{tc}$  do  
13        forall  $(s_2, \phi_2, \uparrow_i^I(\alpha), \rho_2, a, s') \in \Delta$  with  $\nu_2 \models \phi_2, t \in I$  do  
14           $\mathcal{R}_{tc} := \mathcal{R}_{tc} \cup (s, \nu, t, s', \rho_2[\nu_2])$ ;  
15     $\mathcal{R}_{tc} := \text{TransitiveClosureTimed}(\mathcal{R}_{tc})$ ;  
16    if  $\mathcal{R}_{tc} \setminus \text{WRT} = \emptyset$  then  
17      break;  
18 return  $\text{WRT}$ ;
```

Algorithm 11: Add Hole Timed

```
1 Function AddHoleTimed $_i(\mu, HST_i)$ :  
   Result:  $Set = \{ \mu \mid \mu \text{ is a list of states and time elapses} \}$   
2    $Set := \emptyset$ ;  
3    $(s, \nu) := \text{last}(\mu)$ ;  
4   forall  $(i, s, \nu, t, s', \nu') \in HST_i$  do  
5      $\mu' = \text{copy}(\mu)$ ;  
6      $\text{trunc}(\mu')$ ; /*  $\text{trunc}(\mu)$  is defined as  $\text{remove}(\text{last}(\mu))$  */  
7      $\mu'.\text{append}[(i, s, \nu, t, s', \nu'), 0, (s', \nu')]$ ;  
8      $\mu'.\text{NumberOfHoles} := \mu.\text{NumberOfHoles} + 1$ ;  
9      $Set := Set \cup \{\mu'\}$ ;  
10 return  $Set$ ;
```

Algorithm 12: Extend with a pop Timed

```

1 Function
  AddPopTimedi( $\mu, M = (\mathcal{S}, \Delta, s_0, \mathcal{S}_f, \mathcal{X}, n, \Sigma, \Gamma), AHST_i, HST_i, WRT$ ):
    Result:  $Set = \{ \mu \mid \mu \text{ is a list of states and time elapses} \}$ 
2    $Set := \emptyset;$ 
3    $[t_l, (s, \nu)] := \text{last}(\mu);$ 
4    $[t', (i, s_1, \nu_1, t, s_3, \nu_3), t''] := \text{lastHole}_i(\mu);$ 
5    $t_3 :=$  The sum of the time elapses in the list  $\mu$  between  $(s_2, \nu_2)_{R_i}$  and  $(s, \nu);$ 
6   forall  $(i, s_1, \nu_1, t_1, s_2, \nu_2) \in HST_i, (i, s_2, \nu_2, t_2, \alpha, s_3, \nu_3) \in AHST_i,$ 
      $(s, \phi, R, \uparrow_i^I(\alpha), s') \in \Delta$  with  $t = t_1 + t_2, \nu \models \phi$  and  $t_2 + t_3 \in I,$  and
      $(s', R[\nu], t_4, s'', \nu'') \in WRT$  do
7      $\mu' = \text{copy}(\mu);$ 
8      $\text{trunc}(\mu');$ 
9      $\mu'.\text{append}([t_l \oplus t_4, (s'', \nu'')]);$ 
10     $\mu'.\text{replace}([t', (i, s_1, \nu_1, t, s_3, \nu_3), t''],$ 
         $[t', (i, s_1, \nu_1, t_1, s_2, \nu_2), t_2 \oplus t'']);$ 
11     $Set := Set \cup \{\mu'\};$ 
12    if  $t_1 = 0$  and  $(s_1, \nu_1) = (s_2, \nu_2)$  then
13       $\mu'' = \text{copy}(\mu);$ 
14       $\text{trunc}(\mu'');$ 
15       $\mu''.\text{append}([t_l \oplus t_4, (s'', \nu'')]);$ 
16       $\mu''.\text{replace}([t', (i, s_1, \nu_1, t, s_3, \nu_3), t''], (t' \oplus t \oplus t''));$ 
17       $\mu''.\text{NumberOfHoles} = \mu.\text{NumberOfHoles} - 1;$ 
18       $Set := Set \cup \{\mu''\};$ 
19 return  $Set;$ 

```

1075 F Witness Generation for TMPDA

1076 In this section, we focus on the important question of generating a witness for
 1077 an accepting run whenever our fixed-point algorithm guarantees non-emptiness.
 1078 Since we use fixed-point computations to speed up our reachability algorithm,
 1079 finding a witness, i.e., an explicit run witnessing reachability, becomes non-trivial.
 1080 In fact, the difficulty of the witness generation depends on the system under
 1081 consideration : while it is reasonably straight-forward for timed automata with
 1082 no stacks, it is quite non-trivial when we have (multiple) stacks with non-well
 1083 nested behavior.

Algorithm 13: Well-nested Timed Witness Generation

```

1 Function WitnessTimedWR( $s_1, s_2, \nu, M = (\mathcal{S}, \Delta, s_0, \mathcal{S}_f, \mathcal{X}, n, \Sigma, \Gamma), WRT$ ):
   Result: A sequence of transitions for an accepting run
2   if  $s_1 == s_2$  then
3     return  $\epsilon$ ;
4   if  $\exists t = (s, \phi, R, \text{nop}, s') \in \Delta \wedge \nu \models \phi \wedge \nu = R[\nu]$  then
5     return  $t$ ;
6   forall  $s', s'' \in \mathcal{S}$  do
7     if  $((s_1 \neq s') \vee (s'' \neq s_2)) \wedge (s', s'') \in WRT$ 
        $\wedge \exists t = (s_1, \phi, R, \downarrow_i(\alpha), a, s') \in \Delta \wedge$ 
        $\exists t_2 = (s'', \phi', R', \uparrow_i(\alpha), a', s_2) \in \Delta \wedge \nu = R[\nu] = R[\nu'] \wedge \nu \models \phi \wedge \nu \models \phi'$ 
       then
8        $\text{path} = \text{WitnessTimedWR}(s', s'', \nu, M, WRT)$ ;
9       return  $t.\text{path}.t_2$ ;
10  forall  $s \in M.S$  do
11    if  $(s \neq s_1 \vee s \neq s_2) \wedge (s, 0, s_1) \in WRT \wedge (s, 0, s_2) \in WRT$  then
12       $\text{path1} = \text{WitnessTimedWR}(s_1, s, \nu, M, WRT)$ ;
13       $\text{path2} = \text{WitnessTimedWR}(s, s_2, \nu, M, WRT)$ ;
14      return  $\text{path1}.\text{path2}$ ;

```

1084 **0-holes.** We start discussing the witness generation in the case of timed automata.
 1085 As described in the algorithm in section 3, non-emptiness is guaranteed if a final
 1086 state (s_f, ν_f) is reached from the initial state (s_0, ν_0) by computing the transitive
 1087 closure of the transitions. The transitive closure computation results in generating
 1088 a tuple $(s_0, \nu_0, t, s_f, \nu_f) \in WRT$ (Algorithm 10), for some time $0 \leq t \in \mathbb{R}$. Notice
 1089 however that, in the Algorithms 10, we do not keep track of the sequence
 1090 of states that led to the final state, and this is why we need to reconstruct a
 1091 witness. To generate a witness run, we consider a *normal form* for any run in
 1092 the underlying timed automaton, and check for the existence of a witness in the
 1093 normal form. A run is in the normal form if it is a sequence of *time-elapse*, *useful*,
 1094 and *useless* transitions. Time-elapse transitions have already been explained
 1095 earlier. A discrete transition $(s, \nu) \rightarrow (s', \nu')$ is *useful* if $\nu \neq \nu'$, that is, there is
 1096 at least one clock x such that $\nu'(x) = 0$ and $\nu(x) \neq 0$. A discrete transition is
 1097 *useless* if $\nu = \nu'$.

Algorithm 14: Timed Pushdown Automata Witness Generation

```

1 Function Witness( $(s_1, \nu_1), t, (s_2, \nu_2), M = (\mathcal{S}, \Delta, s_0, \mathcal{S}_f, \mathcal{X}, n, \Sigma, \Gamma), \text{WRT})$ :
   Result: A sequence of transitions for an accepting run
2   forall  $t_1 \in [T]$  do
3      $\text{midPath} = \text{Witness}((s_1, \nu_1 + t_1), t - t_1, (s_2, \nu_2), M, \text{WRT})$  Progress Measure 1;
4     if  $\text{midPath} \neq \emptyset$  then
5       return  $t_1 \cdot \text{midPath}$ ;
6   forall  $\delta = (s'', \phi', R', \text{nop}, a', s_2) \in M \cdot \Delta$  do
7     if  $\delta.R'[\nu_1] \neq \nu_1$  and  $\nu_1 \models \delta.\phi'$  then
8        $s_3 = \delta.s_2$ ;
9        $\nu_3 = \delta.R'[\nu_1]$ ;
10       $\text{midPath2} = \text{Witness}((s_3, \nu_3), t, (s_2, \nu_2), M, \text{WRT})$  Progress Measure 2;
11      if  $\text{midPath2} \neq \emptyset$  then
12        return  $\delta \cdot \text{midPath2}$ ;
13   forall  $s \in M.S$  do
14      $\text{path} = \text{WitnessTimedWR}(s_1, s, \nu_1, M, \text{WRT})$  Progress Measure 3;
15     if  $\text{path} \neq \emptyset$  then
16        $\text{midPath3} = \text{Witness}((s, \nu_1), t, (s_2, \nu_2), M, \text{WRT})$ ;
17       if  $\text{midPath3} \neq \emptyset$  then
18         return  $\text{path} \cdot \text{midPath3}$ ;

```

1098 If a tuple $(s_0, \nu_0, t, s_f, \nu_f)$, $t \geq 0$ is generated by Algorithm 10, we know
 1099 that the system is non-empty. Now, we describe an algorithm to generate the
 1100 witness run for obtaining $(s_0, \nu_0, t, s_f, \nu_f)$, by associating a *lexicographic progress*
 1101 *measure* while exploring runs starting from (s_0, ν_0) . Integral time elapses, useful
 1102 transitions and useless transitions are the three entities constituting the progress
 1103 measure, ordered lexicographically.

- 1104 – First we check if it is possible to obtain a witness run of the form $(s_0, \nu_0) \xrightarrow{t_1}$
 1105 $(s, \nu) \xrightarrow{t_2} (s_f, \nu_f)$, where \xrightarrow{t} denotes a sequence of transitions whose total time
 1106 elapse is t . In case $t_1, t_2 > 0$, with $t_1 + t_2 = t$, we can recurse on obtaining
 1107 witnesses to reach (s, ν) from (s_0, ν_0) , and (s_f, ν_f) from (s, ν) , with strictly
 1108 smaller time elapses, guaranteeing progress to termination.
- 1109 – In case $t_1 = 0$ or $t_2 = 0$, we move to the second component of our progress
 1110 measure, namely useful transitions. Assume $t_2 = 0$. Then indeed, there
 1111 is no time elapse in reaching (s_f, ν_f) from (s, ν) , but only a sequence of
 1112 discrete transitions. Let $\#_X(\nu)$ denote the number of non-zero entries in the
 1113 valuation ν . To obtain the witness, we look at a maximal sequence of useful
 1114 transitions from (s, ν) of the form $(s, \nu) \rightarrow (s_1, \nu_1) \rightarrow \dots \rightarrow (s_k, \nu_k)$ such
 1115 that $\#_X(\nu) > \#_X(\nu_1) > \dots > \#_X(\nu_k)$, where $k \leq$ the number of clocks.
 1116 When we reach some (s_i, ν_i) from where we cannot make a useful transition,
 1117 we go for a useless transition. Since there is no time elapse, and no useful
 1118 resets, the clock valuations do not change on discrete transitions. We are left
 1119 with enumerating all the locations to check the reachability to s_f (or to some
 1120 s_j , from where we can again have a maximal sequence of useful transitions).
 1121 Indeed, if (s_f, ν_f) is reachable from (s, ν) with no time elapse, there is a path
 1122 having at most $|\mathcal{X}|$ useful transitions, interleaved with a sequence of useless
 1123 transitions.

1124 Generation of witness for timed automata is given in Algorithm 14. Notice that
 1125 when $\kappa = (s_0, \nu_0, 0, s_f, \nu_f)$, the progress measure is $m(\kappa) = \#_X(\nu_0) - \#_X(\nu_f)$.

1126 If $m(\kappa) = 0$, then $\nu_0 = \nu_f$, and the path takes only useless transitions. In this
 1127 case, we consider the graph with nodes as states (s, ν) , and there is an edge
 1128 from (s_1, ν_1) to (s_2, ν_2) if there is a transition (s_1, φ, R, s_2) such that $\nu_1 \models \varphi$
 1129 and $\nu_1[R] = \nu_2$, that is, for all $x \in R$, $\nu_1(x) = 0$. If $m(\kappa) \neq 0$, then we take at
 1130 least one useful transition. We can check if there exists a transition (s_1, φ, R, s_2)
 1131 such that s_1 is reachable from s_0 , and $\nu_0 \models \varphi, \nu_0[R] \neq \nu_0$, and the tuple
 1132 $\kappa' = (s_2, \nu_0[R], 0, s_f, \nu_f) \in \text{WRT}$. In this case, we have $m(\kappa') < m(\kappa)$ and we can
 1133 conclude by induction.

1134 The case of a timed pushdown system with a single stack is similar to the
 1135 case of timed automata, except for the fact that a discrete transition may
 1136 involve push/pop operations. We use the same progress measures as in the timed
 1137 automaton case, using the notion of runs in normal form.

1138 **Getting Witness from Holes.** We can extend the backtracking algorithm for
 1139 witness generation for MPDA to generate witness for TMPDA without much
 1140 modification. In timed settings we need to take care of the time elapses within a
 1141 hole and an atomic hole segment. When a hole is partitioned to an atomic hole
 1142 segment and a hole, the time must be partitioned satisfying possible atomic hole
 1143 segments and holes along with other constraints.