

# Revisiting Underapproximate Reachability for Multipushdown Systems

S. Akshay<sup>1</sup>, Paul Gastin<sup>2</sup>, Krishna S<sup>1</sup>, and Sparsa Roychowdhury<sup>1</sup>

<sup>1</sup> IIT Bombay, Mumbai, India  
{akshayss,krishnas,sparsa}@cse.iitb.ac.in  
<sup>2</sup> ENS Paris-Saclay, France  
paul.gastin@lsv.fr

**Abstract.** Boolean programs with multiple recursive threads can be captured as pushdown automata with multiple stacks. This model is Turing complete, and hence, one is often interested in analyzing a restricted class which still captures useful behaviors. In this paper, we propose a new class of bounded underapproximations for multipushdown systems, which subsumes most existing classes. We develop an efficient algorithm for solving the under-approximate reachability problem, which is based on efficient fix-point computations. We implement it in our tool **BHIM** and illustrate its applicability by generating a set of relevant benchmarks and examining its performance. As an additional takeaway **BHIM** solves the binary reachability problem in pushdown automata. To show the versatility of our approach, we then extend our algorithm to the timed setting and provide the first implementation that can handle timed multi-pushdown automata with closed guards.

**Keywords:** Multipushdown Systems, Underapproximate Reachability, Timed pushdown automata

## 1 Introduction

The reachability problem for pushdown systems with multiple stacks is known to be undecidable. However, multi-stack pushdown automata (MPDA hereafter) represent a theoretically concise and analytically useful model of multi-threaded recursive programs with shared memory. As a result, several previous works in the literature have proposed different under-approximate classes of behaviors of MPDA that can be analyzed effectively, such as *Round Bounded*, *Scope Bounded*, *Context Bounded* and *Phase Bounded* [1,2,3,4,5,6]. From a practical point of view, these underapproximations has led to efficient tools including, GetaFix [7], SPADE [8]. It has also been argued (e.g., see [9]) that such bounded underapproximations suffice to find several bugs in practice. In many such tools efficient fix-point techniques are used to speed-up computations.

We extend known fix-point based approaches by developing a new algorithm that can handle a larger class of bounded underapproximations than bounded

context and bounded scope for multi-pushdown systems while remaining efficiently implementable. This algorithm works for a new class of underapproximate behaviors called *hole bounded* behaviors, which subsumes context or scope bounded underapproximations, and is orthogonal to phase bounded underapproximations. A “hole” is a maximal sequence of push operations of a fixed stack, interspersed with well-nested sequences of any stack. Thus, in a sequence  $\alpha = \beta\gamma$  where  $\beta = [push_1(push_2push_3pop_3pop_2)push_1(push_3pop_3)]^{10}$  and  $\gamma = push_2push_1pop_2pop_1(pop_1)^{20}$ ,  $\beta$  is a hole wrt stack 1. The suffix  $\gamma$  has 2 holes (the  $push_2$  and the  $push_1$ ). The number of context switches in  $\alpha$  is  $> 50$ , and so is the number of changes in scope, while  $\alpha$  is 3-hole bounded. A ( $k$ )-hole bounded sequence is one such, where, at any point of the computation, the number of holes are bounded (by  $k$ ). We show that the class of hole bounded sequences subsumes most of the previously defined classes of underapproximations and is, in fact, contained in the very generic class of tree-width bounded sequences. This immediately shows decidability of reachability for our class.

Analyzing the more generic class of tree-width bounded sequences is often much more difficult; for instance, building bottom-up tree automata for this purpose does not scale very well as it explores a large (and often useless) state space. Our technique is radically different from using tree automata. Under the hole-bounded assumption, we pre-compute information regarding well-nested sequences and holes using fix-point computations and use them in our algorithm. Using efficient data structures to implement this approach, we develop a tool (BHIM) for Bounded Hole reachability in Multistack pushdown systems.

#### Highlights of BHIM.

- Two significant aspects of the fix-point approach in BHIM are: we efficiently solve the binary reachability problem for pushdown automata. i.e., BHIM computes all pairs of states  $(s, t)$  such that  $t$  is reachable from  $s$  with empty stacks. This allows us to go beyond reachability and handle some liveness questions; (ii) we pre-compute the set of pairs of states that are endpoints of holes. This allows us to greatly limit the search for an accepting run.
- While the fix-point approach solves (binary) reachability efficiently, it does not a priori produce a witness of reachability. We remedy this situation by proposing a backtracking algorithm, which cleverly uses the computations done in the fix-point algorithm, to generate a witness efficiently.
- BHIM is parametrized w.r.t the hole bound: if non-emptiness can be checked or witnessed by a well-nested sequence (this is an easy witness and BHIM looks for easy witnesses first, then gradually increases complexity, if no easy witness is found), then it is sufficient to have the hole bound 0; increasing this complexity measure as required to certify non-emptiness gives an efficient implementation, in the sense that we search for harder witnesses only when no easier witnesses (w.r.t this complexity measure) exist. In all examples as described in the experimental section, a small (less than 4) bound suffices and we expect this to be the case for most practical examples.
- Finally, extend our approach to handle timed multi-stack pushdown systems. This shows the versatility of our approach and also requires us to solve several

83 technical challenges which are specific to the timed setting. Implementing this  
84 approach in **BHIM** makes it, to the best of our knowledge, the first tool that can  
85 analyze timed multi-stack pushdown automata (TMPDA) with closed guards.

86 We analyze the performance of **BHIM** in practice, by considering benchmarks  
87 from the literature, and generating timed variants of some of them. We modeled  
88 two variants of the Bluetooth example [10,8] and **BHIM** was able to detect three  
89 errors (of which it seems only two were already known). Likewise, for an example  
90 of a multiple producer consumer model, **BHIM** could detect bugs by finding  
91 witnesses having just 3 holes, while, it is unlikely that existing tools working  
92 on scope/context bounded underapproximations can handle them as the no. of  
93 switches in scope/context required would exceed 40 to find the bug. In the timed  
94 setting, one of the main challenges faced has been the unavailability of timed  
95 benchmarks; even in the untimed setting, many benchmarks were unavailable  
96 due to their proprietary nature. Nevertheless we tested our tool on 5 other  
97 benchmarks and 3 timed variants whose details, along with their parametric  
98 dependence plots, are given in Supplementary Material [11]. Due to lack of space  
99 proofs and technical details, especially in the timed setting are also in [11].



100 **Related Work.** Among other under-approximations, scope bounded [3] subsumes  
101 context and round bounded underapproximations, and it also paves path for  
102 GetaFix [7], a tool to analyze recursive (and multi-threaded) boolean programs.  
103 As mentioned earlier hole-boundedness strictly subsumes scope boundedness. On  
104 the other hand, GetaFix uses symbolic approaches via BDDs, which is orthogonal  
105 to the improvements made in this paper. Indeed, our next step would be to  
106 build a symbolic version of **BHIM** which extends the hole-bounded approach to  
107 work with symbolic methods. Given that **BHIM** can already handle synthetic  
108 examples with 12-13 holes (see [11]), we expect this to lead to even more drastic  
109 improvements and applicability. For sequential programs, a summary-based  
110 algorithm is used in [7]; summaries are like our well-nested sequences, except that  
111 well-nested sequences admit contexts from different stacks unlike summaries. As  
112 a result, our class of bounded hole behaviors generalizes summaries. Many other  
113 different theoretical results like phase bounded [1], order bounded [12] which gives  
114 interesting underapproximations of MPDA, are subsumed in tree-width bounded  
115 behaviors, but they do not seem to have practical implementations. Adding  
116 real-time information to pushdown automata by using clocks or timed stacks has  
117 been considered, both in the discrete and dense-timed settings. Recently, there  
118 has been a flurry of theoretical results in the topic [13,14,15,16,17]. However,  
119 to the best of our knowledge none of these algorithms have been successfully  
120 implemented (except [17] which implements a tree-automata based technique  
121 for single-stack timed systems) for multi-stack systems. One reason is that these  
122 algorithms do not employ scalable fix-point based techniques, but instead depend  
123 on region automaton-based search or tree automata-based search techniques.

## 124 2 Underapproximations in MPDA

125 A multi-stack pushdown automaton (MPDA) is a tuple  $M = (\mathcal{S}, \Delta, s_0, \mathcal{S}_f, n, \Sigma, \Gamma)$   
126 where,  $\mathcal{S}$  is a finite non-empty set of locations,  $\Delta$  is a finite set of transitions,

127  $s_0 \in \mathcal{S}$  is the initial location,  $\mathcal{S}_f \subseteq \mathcal{S}$  is a set of final locations,  $n \in \mathbb{N}$  is the  
 128 number of stacks,  $\Sigma$  is a finite input alphabet, and  $\Gamma$  is a finite stack alphabet  
 129 which contains  $\perp$ . A transition  $t \in \Delta$  can be represented as a tuple  $(s, \text{op}, a, s')$ ,  
 130 where,  $s, s' \in \mathcal{S}$  are respectively, the source and destination locations of the  
 131 transition  $t$ ,  $a \in \Sigma$  is the label of the transition, and  $\text{op}$  is one of the following  
 132 operations (1) **nop**, or no stack operation, (2)  $(\downarrow_i \alpha)$  which pushes  $\alpha \in \Gamma$  onto  
 133 stack  $i \in \{1, 2, \dots, n\}$ , (3)  $(\uparrow_i \alpha)$  which pops stack  $i$  if the top of stack  $i$  is  $\alpha \in \Gamma$ .

134 For a transition  $t = (s, \text{op}, a, s')$  we write  $\text{src}(t) = s$ ,  $\text{tgt}(t) = s'$  and  $\text{op}(t) = \text{op}$ .  
 135 At the moment we ignore the action label  $a$  but this will be useful later when we  
 136 go beyond reachability to model checking. A *configuration* of the MPDA is a tuple  
 137  $(s, \lambda_1, \lambda_2, \dots, \lambda_n)$  such that,  $s \in \mathcal{S}$  is the current location and  $\lambda_i \in \Gamma^*$  represents  
 138 the current content of  $i^{\text{th}}$  stack. The semantics of the MPDA is defined as follows:  
 139 a run is accepting if it starts from the initial state and reaches a final state with  
 140 all stacks empty. The language accepted by a MPDA is defined as the set of words  
 141 generated by the accepting runs of the MPDA. Since the reachability problem for  
 142 MPDA is Turing complete, we consider under-approximate reachability.

143 A sequence of transitions is called **complete** if each push in that sequence  
 144 has a matching pop and vice versa. A **well-nested** sequence denoted  $ws$  is  
 145 defined inductively as follows: a possibly empty sequence of **nop**-transitions is  
 146  $ws$ , and so is the sequence  $t ws t'$  where  $\text{op}(t) = (\downarrow_i \alpha)$  and  $\text{op}(t') = (\uparrow_i \alpha)$  are a  
 147 matching pair of push and pop operations of stack  $i$ . Finally the concatenation  
 148 of two well-nested sequences is a well-nested sequence, i.e., they are closed under  
 149 concatenation. The set of all well-nested sequences defined by an MPDA is  
 150 denoted **WS**. If we visualize this by drawing edges between pushes and their  
 151 corresponding pops, well-nested sequences have no crossing edges, as in   
 152 and , where we have two stacks, depicted with red and violet edges. We  
 153 emphasize that a well-nested sequence can have well-nested edges from any stack.  
 154 In a sequence  $\sigma$ , a push (pop) is called a **pending** push (pop) if its matching  
 155 pop (push) is not in the same sequence  $\sigma$ .

156 **Bounded Underapproximations.** As mentioned in the introduction, different  
 157 bounded under-approximations have been considered in the literature to get  
 158 around the Turing completeness of MPDA. During a computation, a context is a  
 159 sequence of transitions where only one stack or no stack is used. In *context bounded*  
 160 computations the number of contexts are bounded [18]. A *round* is a sequence  
 161 of (possibly empty) contexts for stacks  $1, 2, \dots, n$ . *Round bounded* computations  
 162 restrict the total number of rounds allowed [2,16,17]. *Scope bounded* computations  
 163 generalize bounded context computations. Here, the context changes within any  
 164 push and its corresponding pop is bounded [2,5,6]. A *phase* is a contiguous  
 165 sequence of transitions in a computation, where we restrict pop to only one stack,  
 166 but there are no restrictions on the pushes [1]. A phase bounded computation is  
 167 one where the number of phase changes is bounded.

168 **Tree-width.** A generic way of looking at them is to consider classes which have a  
 169 bound on the tree-width [19]. In fact, the notions of split-width/cliue-width/tree-  
 170 width of communicating finite state machines/timed push down systems has been  
 171 explored in [20], [21]. The behaviors of the underlying system are then represented

as graphs. It has been shown in these references that if the family of graphs arising from the behaviours of the underlying system (say  $S$ ) have a bounded tree-width, then the reachability problem is decidable for  $S$  via, tree-automata. However, this does not immediately give rise to an efficient implementation. The tree-automata approach usually gives non-deterministic or bottom-up tree automata, which when implemented in practice (see [17]) tend to blow up in size and explore a large and useless space. Hence there is a need for efficient algorithms, which exist for more specific underapproximations such as context-bounded (leading to fix-point algorithms and their practical implementations [7]).


## 2.1 A new class of under-approximations

Our goal is to bridge the gap between having practically efficient algorithms and handling more expressive classes of under-approximations for reachability of multi-stack pushdown systems. To do so, we define a bounded approximation which is expressive enough to cover previously defined practically interesting classes (such as context bounded etc), while at the same time allowing efficient decidable reachability tests, as we will see in the next section.

**Definition 1.** (*Holes*). Let  $\sigma$  be complete sequence of transitions, of length  $n$  in a MPDA, and let  $ws$  be a (possibly empty) well-nested sequence.

- A **hole** of stack  $i$  is a maximal factor of  $\sigma$  of the form  $(\downarrow_i ws)^+$ , where  $ws \in \text{WS}$ . The maximality of the hole of stack  $i$  follows from the fact that any possible extension ceases to be a hole of stack  $i$ ; that is, the only possible events following a maximal hole of stack  $i$  are a push  $\downarrow_j$  of some stack  $j \neq i$ , or a pop of some stack  $j \neq i$ . In general, whenever we speak about a hole, the underlying stack is clear.
- A push  $\downarrow_i$  in a hole (of stack  $i$ ) is called a pending push at (i.e., just before) a position  $x \leq n$ , if its matching pop occurs in  $\sigma$  at a position  $z > x$ .
- A hole (of stack  $i$ ) is said to be **open** at a position  $x \leq n$ , if there is a pending push  $\downarrow_i$  of the hole at  $x$ . Let  $\#_x(\text{hole})$  denote the number of open holes at position  $x$ . The **hole bound** of  $\sigma$  is defined as  $\max_{1 \leq x \leq |\sigma|} \#_x(\text{hole})$ .
- A hole segment of stack  $i$  is a prefix of a hole of stack  $i$ , ending in a  $ws$ , while an atomic hole segment of stack  $i$  is just the segment of the form  $\downarrow_i ws$ .

As an example, consider the sequence  $\sigma$  in Figure 1 of transitions of a MPDA having stacks 1,2 (denoted respectively red and blue). We use superscripts for each push, pop of each stack to distinguish the  $i$ th push,  $j$ th pop and so on of each stack. There are two holes of stack 1 (red stack) denoted by the red patches,



**Fig. 1.** A run  $\sigma$  with 2 holes (2 red patches) of the red stack and 1 hole (one blue patch) of the blue stack.

and one hole of stack 2 (blue stack) denoted by the blue patch. The subsequence  $\downarrow_1^1 \downarrow_1^2 ws_2$  of the first hole is not a maximal factor, since it can be extended by  $\downarrow_1^3 ws_3$  in the run  $\sigma$ , extending the hole. Consider the position in  $\sigma$  marked with  $\downarrow_2^1$ . At this position, there is an open hole of the red stack (the first red patch), and there is an open hole of the blue stack (the blue patch). Likewise, at the position  $\uparrow_1^5$ , there are 2 open holes of the red stack (2 red patches) and one open hole of the blue stack 2 (the blue patch). The hole bound of  $\sigma$  is 3. The green patch consisting of  $\uparrow_1^3, \uparrow_1^2$  and  $ws_5$  is a pop-hole of stack 1. Likewise, the pops  $\uparrow_2^2, \uparrow_1^5, \uparrow_2^1$  are all pop-holes (of length 1) of stacks 2,1,2 respectively.

**Definition 2.** (HOLE BOUNDED REACHABILITY PROBLEM) *Given a MPDA and  $K \in \mathbb{N}$ , the  $K$ -hole bounded reachability problem is the following: Does there exist a  $K$ -hole bounded accepting run of the MPDA?*

**Proposition 1.** *The tree-width of  $K$ -hole bounded MPDA behaviors is at most  $(2K + 3)$ .*

A detailed proof of this Proposition is given in Appendix ???. Once we have this, from [19][16][17], decidability and complexity follow immediately. Thus,

**Corollary 1.** *The  $K$ -hole bounded reachability problem for MPDA is decidable in  $\mathcal{O}(|\mathcal{M}|^{2K+3})$  where,  $\mathcal{M}$  is the size of the underlying MPDA.*

Next, we turn to the expressiveness of this class wrt to the classical underapproximations of MPDA: first, the **hole** bounded class strictly subsumes **scope** bounded which already subsumes **context** bounded and **round** bounded classes. Also **hole** bounded MPDA and **phase** bounded MPDA are orthogonal.

**Proposition 2.** *Consider a MPDA  $M$ . For any  $K$ , let  $L_K$  denote a set of sequences accepted by  $M$  which have number of rounds or number of contexts or scope bounded by  $K$ . Then there exists  $K' \leq K$  such that  $L_K$  is  $K'$  hole bounded. Moreover, there exist languages which are  $K$  hole bounded for some constant  $K$ , which are not  $K'$  round or context or scope bounded for any  $K'$ . Finally, there exists a language which is accepted by phase bounded MPDA but not accepted by hole bounded MPDA and vice versa.*

*Proof.* We first recall that if a language  $L$  is  $K$ -round, or  $K$ -context bounded, then it is also  $K'$ -scope bounded for some  $K' \leq K$  [5,2]. Hence, we only show that scope bounded systems are subsumed by hole bounded systems.

Let  $L$  be a  $K$ -scope bounded language, and let  $M$  be a MPDA accepting  $L$ . Consider a run  $\rho$  of  $w \in L$  in  $M$ . Assume that at any point  $i$  in the run  $\rho$ ,  $\#_i(\text{holes}) = k'$ , and towards a contradiction, let,  $k' > K$ . Consider the leftmost open hole in  $\rho$  which has a pending push  $\downarrow^p$  whose pop  $\uparrow^p$  is to the right of  $i$ . Since  $k' > K$  is the number of open holes at  $i$ , there are at least  $k' > K$  context changes in between  $\downarrow^p$  and  $\uparrow^p$ . This contradicts the  $K$ -scope bounded assumption, and hence  $k' \leq K$ .

To show the strict containment, consider the visibly pushdown language [22] given by  $L^{bh} = \{a^n b^n (a^{p_1} c^{p_1+1} b^{p'_1} d^{p'_1+1} \dots a^{p_n} c^{p_n+1} b^{p'_n} d^{p'_n+1}) \mid n, p_1, p'_1, \dots, p_n, p'_n \in$

248  $\mathbb{N}\}$ . A possible word  $w \in L^{bh}$  is  $a^3b^3 a^2c^3b^2d^3 a^2c^3bd^2 ac^2bd^2$  with  $a, b$  representing  
 249 push in stack 1,2 respectively and  $c, d$  representing the corresponding matching  
 250 pop from stack 1,2. A run  $\rho$  accepting the word  $w \in L^{bh}$  will start with a sequence  
 251 of pushes of stack 1 followed by another sequence of pushes of stack 2. Note that,  
 252 the number of the pushes  $n$  is same in both stacks. Then there is a group  $G$   
 253 consisting of a well-nested sequence of stack 1 (equal  $a$  and  $c$ ) followed by a pop  
 254 of the stack 1 (an extra  $c$ ), another well-nested sequence of stack 2 (equal  $b$  and  
 255  $d$ ) and a pop of the stack 2 (an extra  $d$ ), repeated  $n$  times. From the definition  
 256 of the `hole`, the total number of holes required in  $G$  is 0. But, we need 1 hole for  
 257 the sequence of  $a$ 's and another for the sequence of  $b$ 's at the beginning of the  
 258 run, which creates at most 2 holes during the run. Thus, the hole bound for any  
 259 accepting run  $\rho$  is 2, and the language  $L^{bh}$  is 2-hole bounded.

However,  $L^{bh}$  is not  $k$ -scope bounded for any  $k$ . Indeed, for each  $m \geq 1$ ,  
 consider the word  $w_m = a^mb^m(ac^2bd^2)^m \in L^{bh}$ . It is easy to see that  $w_m$  is  
 2m-scope bounded (the matching  $c, d$  of each  $a, b$  happens  $2m$  context switches  
 later) but not  $k$ -scope bounded for  $k < 2m$ . It can be seen that  $L^{bh}$  is not  $k$ -phase  
 bounded either. Finally,  $L' = \{(ab)^nc^nd^n \mid n \in \mathbb{N}\}$  with  $a, b$  and  $c, d$  respectively  
 being push and pop of stack 1,2 is not hole-bounded but 2-phase bounded.  $\square$

### 260 3 A Fix-point Algorithm for Hole Bounded Reachability

261 In the previous section, we showed that hole-bounded underapproximations are  
 262 a decidable subclass for reachability, by showing that this class has a bounded  
 263 tree-width. However, as explained in the introduction, this does not immediately  
 264 give a fix-point based algorithm, which has been shown to be much more efficient  
 265 for other more restricted sub-classes, e.g., context-bounded. In this section, we  
 266 provide such a fix-point based algorithm for the hole-bounded class and explain its  
 267 advantages. Later we discuss its versatility by showing extensions and evaluating  
 268 its performance on a suite of benchmarks.

269 We describe the algorithm in two steps: first we give a simple fix-point based  
 270 algorithm for the problem of 0-hole or *well-nested reachability*, i.e, reachability  
 271 by a well-nested sequence without any holes. For the 0-hole case, our algorithm  
 272 computes the *reachability relation*, also called the *binary reachability problem* [23].  
 273 That is, we accept all pairs of states  $(s, s')$  such that there is a well-nested run  
 274 from  $s$  with empty stack to  $s'$  with empty stack. Subsequently, we combine this  
 275 binary reachability for well-nested sequences with an efficient graph search to  
 276 obtain an algorithm for  $K$ -hole bounded reachability.

277 **Binary well-nested reachability for MPDA.** Note that single stack PDA are  
 278 a special case, since all runs are indeed well-nested.

- 279 1. **Transitive Closure:** Let  $\mathcal{R}$  be the set of tuples of the form  $(s_i, s_j)$  representing  
 280 that state  $s_j$  is reachable from state  $s_i$  via a `nop` discrete transition. Such a  
 281 sequence from  $s_i$  to  $s_j$  is trivially *well-nested*. We take the **TransitiveClosure**  
 282 of  $\mathcal{R}$  using Floyd-Warshall algorithm [24]. The resulting set  $\mathcal{R}_c$  of tuples  
 283 answers the binary reachability for finite state automata (no stacks).

---

**Algorithm 1:** Algorithm for Emptiness Checking of hole bounded MPDA

---

```

1 Function IsEmpty( $M = (\mathcal{S}, \Delta, s_0, S_f, n, \Sigma, \Gamma), K$ ):
   Result: True or False
2  $WR := \text{WellNestedReach}(M)$ ;  $\backslash \backslash$  Solves binary reachability for pushdown system
3 if some  $(s_0, s_1) \in WR$  with  $s_1 \in S_f$  then
4   return False;
5 forall  $i \in [n]$  do
6    $AHS_i := \emptyset$ ;  $Set_i := \emptyset$ ;
7   forall  $(s, \downarrow_i(\alpha), a, s_1) \in \Delta$  and  $(s_1, s') \in WR$  do
8      $AHS_i := AHS_i \cup \{(i, s, \alpha, s')\}$ ;  $Set_i := Set_i \cup \{(s, s')\}$ ;
9    $HS_i := \{(i, s, s') \mid (s, s') \in \text{TransitiveClosure}(Set_i)\}$ ;
10   $\mu := [s_0]$ ;  $\mu.\text{NumberOfHoles} := 0$ ;
11   $SetOfLists_{new} := \{\mu\}$ ;  $SetOfLists := \emptyset$ ;
12  do
13     $SetOfLists := SetOfLists \cup SetOfLists_{new}$ ;
14     $SetOfLists_{todo} := SetOfLists_{new}$ ;  $SetOfLists_{new} := \emptyset$ ;
15    forall  $\mu' \in SetOfLists_{todo}$  do
16      if  $\mu'.\text{NumberOfHoles} < K$  then
17        forall  $i \in [n]$  do
18           $\backslash \backslash$  Add hole for stack i
19           $SetOfLists_h := \text{AddHole}_i(\mu', HS_i) \setminus SetOfLists$ ;
20           $SetOfLists_{new} := SetOfLists_{new} \cup SetOfLists_h$ ;
21        if  $\mu'.\text{NumberOfHoles} > 0$  then
22          forall  $i \in [n]$  do
23             $\backslash \backslash$  Add pop for stack i
24             $SetOfLists_p := \text{AddPop}_i(\mu', M, AHS_i, HS_i, WR) \setminus SetOfLists$ ;
25             $SetOfLists_{new} := SetOfLists_{new} \cup SetOfLists_p$ ;
26            forall  $\mu_3 \in SetOfLists_p$  do
27              if  $\mu_3.\text{last} \in S_f$  and  $\mu_3.\text{NumberOfHoles} = 0$  then
28                return False;  $\backslash \backslash$  If reached destination state
29  while  $SetOfLists_{new} \neq \emptyset$ ;
30  return True;

```

---

284 **2. Push-Pop Closure:** For stack operations, consider a push transition on  
285 some stack (say stack  $i$ ) of symbol  $\gamma$ , enabled from a state  $s_1$ , reaching state  
286  $s_2$ . If there is a matching pop transition from a state  $s_3$  to  $s_4$ , which pops the  
287 same stack symbol  $\gamma$  from the stack  $i$  and if we have  $(s_2, s_3) \in \mathcal{R}_c$ , then we can  
288 add the tuple  $(s_1, s_4)$  to  $\mathcal{R}_c$ . The function `WellNestedReach` (Algorithm ??,  
289 Appendix ??) repeats this process and the transitive closure described above  
290 until a fix-point is reached. Let us denote the resulting set of tuples by  $WR$ .  
291 Thus, we have

292 **Lemma 1.**  $(s_1, s_2) \in WR$  iff  $\exists$  a well-nested run in the MPDA from  $s_1$  to  $s_2$ .

293 **Beyond well-nested reachability.** A naive algorithm for  $K$ -hole bounded  
294 reachability for  $K > 0$  is to start from the initial state  $s_0$ , and do a Breadth  
295 First Search (BFS), nondeterministically choosing between extending with a  
296 well-nested segment, creating hole segments (with a pending push) and closing  
297 hole segments (using pops). We accept when there are no open hole segments  
298 and reach a final state; this gives an exponential time algorithm. Given the  
299 exponential dependence on the hole-bound  $K$  (Corollary 1), this exponential  
300 blowup is unavoidable in the worst case, but we can do much better in practice.  
301 In particular, the naive algorithm makes arbitrary non-deterministic choices  
302 resulting in a blind exploration of the BFS tree.



303 In this section, we use the binary well-nested reachability algorithm as an  
 304 efficient subroutine to limit the search in BFS to its reachable part (note that  
 305 this is quite different from DFS as well since we do not just go down one path).  
 306 The crux is that at any point, we create a new hole for stack  $i$ , *only* when (i)  
 307 we know that we cannot reach the final state without creating this hole and (ii)  
 308 we know that we can close all such holes which have been created. Checking (i)  
 309 is easy, since we just use the WR relation for this. Checking (ii) blindly would  
 310 correspond to doing a DFS; however, we precompute this information and simply  
 311 look it up, resulting in a constant time operation after the precomputation.

312 **Precomputing hole information.** Recall that a *hole* of stack  $i$  is a maximal  
 313 sequence of the form  $(\downarrow_i ws)^+$ , where  $ws$  is a well-nested sequence and  $\downarrow_i$   
 314 represents a push of stack  $i$ . A *hole segment* of stack  $i$  is a prefix of a hole  
 315 of stack  $i$ , ending in a  $ws$ , while an *atomic hole segment* of stack  $i$  is just the  
 316 segment of the form  $\downarrow_i ws$ . A *hole-segment* of stack  $i$  which starts from state  $s$   
 317 in the MPDA and ends in state  $s'$ , can be represented by the triple  $(i, s, s')$ , that  
 318 we call a *hole triple*. We compute the set  $HS_i$  of all hole triples  $(i, s, s')$  such that  
 319 starting at  $s$ , there is a hole segment of stack  $i$  which ends at state  $s'$ , as detailed  
 320 in lines (5-9) of Algorithm 1. In doing so, we also compute the set  $AHS_i$  of all  
 321 atomic hole segments of stack  $i$  and store them as tuples of the form  $(i, s_p, \alpha, s_q)$   
 322 such that  $s_p$  and  $s_q$  are the MPDA states respectively at the left and right end  
 323 points of an atomic hole segment of stack  $i$ , and  $\alpha$  is the symbol pushed on stack  
 324  $i$  ( $s_p \xrightarrow{\downarrow_i(\alpha)ws} s_q$ ).

325 **A guided BFS exploration.** We start with a list  $\mu_0 = [s_0]$  consisting of  
 326 the initial state and construct a BFS exploration tree whose nodes are lists of  
 327 bounded length. A list is a sequence of states and hole triples representing a  
 328  $K$ -hole bounded run in a concise form. If  $H_i$  represents a hole triple for stack  $i$ ,  
 329 then a list is a sequence of the form  $[s, H_i, H_j, H_k, H_i, \dots, H_\ell, s']$ . The simplest  
 330 kind of list is a single state  $s$ . For example, a list with 3 holes of stacks  $i, j, k$  is  
 331  $\mu = [s_0, (i, s, s'), (j, r, r'), (k, t, t'), t']$ . The hole triples (in red) denote open holes  
 332 in the list. The maximum number of open holes in a list is bounded, making the  
 333 length of the list also bounded. Let  $\text{last}(\mu)$  represent the last element of the list  
 334  $\mu$ . This is always a state. For a node  $v$  storing list  $\mu$  in the BFS tree, if  $v_1, \dots, v_k$   
 335 are its children, then the corresponding lists  $\mu_1, \dots, \mu_k$  are obtained by extending  
 336 the list  $\mu$  by one of the following operations:

- 337 1. **Extend  $\mu$  with a hole.** Assume there is a hole of some stack  $i$ , which starts  
 338 at  $\text{last}(\mu) = s$ , and ends at  $s'$ . If the list at the parent node  $v$  is  $\mu = [\dots, s]$ ,  
 339 then for all  $(i, s, s') \in HS_i$ , we obtain the list  $\text{trunc}(\mu) \cdot \text{append}[(i, s, s'), s']$   
 340 at the child node (i.e., we remove the last element  $s$  of  $\mu$ , then append to  
 341 this list the hole triple  $(i, s, s')$ , followed by  $s'$ ). Algorithm ?? in Appendix  
 342 describes this operation in more detail.
- 343 2. **Extend  $\mu$  with a pop.** Suppose there is a transition  $t = (s_k, \uparrow_i(\alpha), a, s'_k)$   
 344 from  $\text{last}(\mu) = s_k$ , where  $\mu$  is of the form  $[s_0, \dots, (h, u, v), (i, s, s'), (j, t, t') \dots, s_k]$ ,  
 345 such that there is no hole triple of stack  $i$  after  $(i, s, s')$ , we extend the run by  
 346 matching this pop (with its push). However, to obtain the last pending push

of stack  $i$  corresponding to this hole, just  $HS_i$  information is not enough since we also need to match the stack content. Instead, we check if we can split the hole  $(i, s, s')$  into (1) a hole triple  $(i, s, s_a) \in HS_i$ , and (2) a tuple  $(i, s_a, \alpha, s') \in AHS_i$ . If both (1) and (2) are possible, then the pop transition  $t$  corresponds to the last pending push of the hole  $(i, s, s')$ .  $t$  indeed matches the pending push recorded in the atomic hole  $(i, s_a, \alpha, s')$  in  $\mu$ , enabling the firing of transition  $t$  from the state  $s_k$ , reaching  $s'_k$ . In this case, we add the child node with the list  $\mu'$  obtained from  $\mu$  as follows. We replace (i)  $s_k$  with  $s'_k$ , and (ii)  $(i, s, s')$  with  $(i, s, s_a)$ , respectively signifying firing of the transition  $t$  and the “shrinking” of the hole, by shifting the end point of the hole segment to the left. When we obtain the hole triple  $(i, s, s)$  (the start and end points of the hole segment coincide), we may have uncovered the last pending push and thereby “closed” the hole segment completely. At this point, we may choose to remove  $(i, s, s)$  from the list, obtaining  $[s_0, \dots, (h, u, v), (j, t, t') \dots, s'_k]$ . For every such  $\mu' = [s_0, \dots, (h, u, v), (i, s, s_a), (j, t, t'), \dots, s'_k]$  and all  $(s'_k, s_m) \in WS$  we also extend  $\mu'$  to  $\mu'' = [s_0, \dots, (h, u, v), (i, s, s_a), (j, t, t'), \dots, s_m]$ . Notice that the size of the list in the child node obtained on a pop, is either the same as the list in the parent, or is smaller. The details are in Algorithm ??.

The number of lists is bounded since the number of states and the length of the lists are bounded. The BFS exploration tree will thus terminate. Combining the above steps gives us Algorithm 1, whose correctness gives us:

**Theorem 1.** *Given a MPDA and a positive integer  $K$ , Algorithm 1 always terminates and answers “false” iff there exists a  $K$ -hole bounded accepting run of the MPDA.*

**Complexity of the Algorithm.** The maximum number of states of the system is  $|\mathcal{S}|$ . The time complexity of transitive closure is  $\mathcal{O}(|\mathcal{S}|^3)$ , using a Floyd-Warshall implementation. The time complexity of Algorithm ??, which uses the transitive closure, is  $\mathcal{O}(|\mathcal{S}|^5) + \mathcal{O}(|\mathcal{S}|^2 \times (|\Delta| \times |\mathcal{S}|))$ . To compute  $AHS$  for  $n$  stacks the time complexity is  $\mathcal{O}(n \times |\Delta| \times |\mathcal{S}|^2)$  and to compute  $HS$  for  $n$  stacks the complexity is  $\mathcal{O}(n \times |\mathcal{S}|^2)$ . For multistack systems, each list keeps track of (i) the number of hole segments ( $\leq K$ ), and (ii) information pertaining to holes (start, end points of holes, and which stack the hole corresponds to). In the worst case, this will be  $(2K + 2)$  possible states in a list, as we are keeping the states at the start and end points of all the hole segments and a stack per hole. So, there are  $\leq |\mathcal{S}|^{2K+3} \times n^{K+1}$  lists. In the worst case, when there is no  $K$ -hole bounded run, we may end up generating all possible lists for a given bound  $K$  on the hole segments. The time complexity is thus bounded above by  $\mathcal{O}(|\mathcal{S}|^{2K+3} \times n^{K+1} + |\mathcal{S}|^5 + |\mathcal{S}|^3 \times |\Delta|)$ .

**Beyond Reachability.** We can solve the usual safety questions in the (bounded-hole) underapproximate setting, by checking for underapproximate reachability on the product of the given system with the complement of the safe set. Given the way Algorithm 1 is designed, the fix-point algorithm allows us to go beyond reachability. In particular, we can solve several (increasingly difficult) variants of the repeated reachability problem, without much modification.

Consider the question : For a given state  $s$  and MPDA, does there exist a run  $\rho$  starting from  $s_0$  which visits  $s$  infinitely often? This is decidable if we can

392 decompose  $\rho$  into a finite prefix  $\rho_1$  and an infinite suffix  $\rho_2$  s.t. (1) Both  $\rho_1, \rho_2$   
 393 are well-nested, or (2)  $\rho_1$  is  $K$ -hole bounded complete (all stacks empty), and  $\rho_2$   
 394 is well-nested, or (3)  $\rho_1$  is  $K$ -hole bounded, and  $\rho_2 = (\rho_3)^\omega$ , where  $\rho_3$  is  $K$ -hole  
 395 bounded. It is easy to see that (1) is solved by two calls to `WellNestedReach` and  
 396 choosing non-empty runs. (2) is solved by a call to Algorithm 1, modified so that  
 397 we reach  $s$ , and then calling `WellNestedReach`. Lastly, to solve (3), first modify  
 398 Algorithm 1 to check reachability to  $s$  with possibly non-empty stacks. Then run  
 399 the modified algorithm twice : first start from  $s_0$  and reach  $s$ ; second start from  
 400  $s$  and reach  $s$  again.

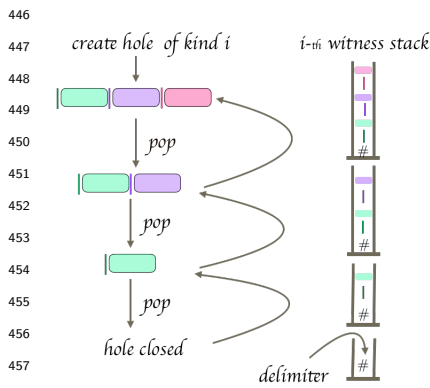
## 401 4 Generating a Witness

402 We next focus on the question of generating a witness for an accepting run  
 403 when our algorithm guarantees non-emptiness. This question is important to  
 404 address from the point of view of applicability: if our goal is to see if bad states  
 405 are reachable, i.e., non-emptiness corresponds to presence of a bug, the witness  
 406 run gives the trace of how the bug came about and hence points to what can  
 407 be done to fix it (e.g., designing a controller). We remark that this question is  
 408 difficult in general. While there are naive algorithms which can explore for the  
 409 witness (thus also solving reachability), these do not use fix-point techniques and  
 410 hence are not efficient. On the other hand, since we use fix-point computations  
 411 to speed up our reachability algorithm, finding a witness, i.e., an explicit run  
 412 witnessing reachability, becomes non-trivial. Generation of a witness in the case of  
 413 well-nested runs is simpler than the case when the run has holes, and requires us  
 414 to “unroll” pairs  $(s_0, s_f) \in \mathbf{WR}$  recursively and generate the sequence of transitions  
 415 responsible for  $(s_0, s_f)$ , as detailed in Algorithm ?? .

416 **Getting Witnesses from Holes.** Now we move on to the more complicated  
 417 case of behaviours having holes. Recall that in BFS exploration we start from  
 418 the states reachable from  $s_0$  by well-nested sequences, and explore subsequent  
 419 states obtained either from (i) a hole creation, or (ii) a pop operation on a stack.  
 420 Proceeding in this manner, if we reach a final configuration (say  $s_f$ ), with all  
 421 holes closed (which implies empty stacks), then we declare non-emptiness. To  
 422 generate a witness, we start from the final state  $s_f$  reachable in the run (a leaf  
 423 node in the BFS exploration tree) and *backtrack* on the BFS exploration tree  
 424 till we reach the initial state  $s_0$ . This results in generating a witness run in the  
 425 reverse, from the right to the left.

426 • Assume that the current node of the BFS tree was obtained using a pop  
 427 operation. There are two possibilities to consider here (see below) depending on  
 428 whether this pop operation closed or shrunk some hole. Recall that each hole  
 429 has a left end point and a right end point and is of a specific stack  $i$ , depending  
 430 on the pending pushes  $\downarrow_i$  it has. So, if the MPDA has  $k$  stacks, then a list in the  
 431 exploration tree can have  $k$  kinds of holes. The witness algorithm uses  $k$  stacks  
 432 called *witness stacks* to correctly implement the backtracking procedure, to deal  
 433 with  $k$  kinds of holes. Witness stacks should not be confused with the stacks of  
 434 the MPDA.

435 • Assume that the current pop operation is closing a hole  $| \text{ } \text{ } \text{ } |$  of kind  
 436  $i$  as in Figure 2. This hole consists of the atomic holes  $| \text{ } \text{ } |$ ,  $| \text{ } \text{ } |$  and  $| \text{ } \text{ } |$ . The  
 437 atomic hole  $| \text{ } \text{ } |$  consists of the push  $|$  and the well-nested sequence  $| \text{ } \text{ } |$  (same  
 438 for the other two atomic holes). Searching among possible push transitions, we  
 439 identify the matching push  $|$  associated with the current pop, resulting in closing  
 440 the hole. On backtracking, this leads to a parent node with the atomic hole  $| \text{ } \text{ } |$   
 441 having as left end point, the push  $|$ , and the right end point as the target of  
 442 the  $ws$   $| \text{ } \text{ } |$ . We push onto the witness stack  $i$ , a barrier (a delimiter symbol  $\#$ )  
 443 followed by the matching push transition  $|$  and then the  $ws$ ,  $| \text{ } \text{ } |$ . The barrier  
 444 segregates the contents of the witness stack when we have two pop transitions of  
 445 the same stack in the reverse run, closing/shrinking two different holes.



459 **Fig. 2.** Backtracking to spit out  
 460 the hole  $| \text{ } \text{ } |$  in reverse. The  
 461 transitions of the atomic hole  $| \text{ } \text{ } |$   
 462 are first written in the reverse  
 463 order, followed by those of  $| \text{ } \text{ } |$   
 464 in reverse, and then of  $| \text{ } \text{ } |$  in reverse.

465 hole segments which constituted this hole. Notice that when we finish processing  
 466 a hole of kind  $i$ , then the witness stack  $i$  has the hole reversed inside it, followed  
 467 by a barrier. The next hole of the same kind  $i$  will be treated in the same manner.

468 • If the current node of the BFS tree is obtained by creating a hole of kind  $i$   
 469 in the fix-point algorithm, then we pop the contents of witness stack  $i$  till we  
 470 reach a barrier. This spits out the atomic hole segments of the hole from the  
 471 right to the left, giving us a sequence of push transitions, and the respective  $ws$   
 472 in between. The transitions constituting the  $ws$  are retrieved using Algorithm ??  
 473 and added. Notice that popping the witness stack  $i$  till a barrier spits out the  
 474 sequence of transitions in the correct reverse order while backtracking.

## 475 5 Adding Time to Multi-pushdown systems

476 In this section, we briefly describe how the algorithms described in section 3  
 477 can be extended to work in the timed setting. Due to lack of space, we focus

478 on some of the significant challenges and advances, leaving the formal details  
 479 and algorithms to the supplement [11]. A TMPDA extends a MPDA with clock  
 480 variables. Transitions check constraints which are conjunctions/disjunctions of  
 481 constraints (called closed guards in the literature) of the form  $x \leq c$  or  $x \geq c$  for  
 482  $c \in \mathbb{N}$  and  $x$  any clock. Symbols pushed on stacks “age” with time elapse. A pop  
 483 is successful only when the age of the symbol lies within a certain interval. The  
 484 acceptance condition is as in the case of MPDA.

485 The first main challenge in adapting the algorithms in section 3 to the timed  
 486 setting was to take care of all possible time elapses along with the operations  
 487 defined in Algorithm 1. The usage of closed guards in TMPDA means that it  
 488 suffices to explore all runs with integral time elapses (for a proof see e.g., Lemma  
 489 4.1 in [16]). Thus configurations are pairs of states with valuations that are vectors  
 490 of non-negative integers, each of which is bounded by the maximal constant in  
 491 the system. Now, to check reachability we need to extend all the precomputations  
 492 (transitive closure, well-nested reachability, as well as atomic and non-atomic hole  
 493 segments) with the time elapse information. To do this, we use a weighted version  
 494 of the Floyd-Warshall algorithm by storing time elapses during precomputations.  
 495 This allows us to use this precomputed *timed* well-nested reachability information  
 496 while performing the BFS tree exploration, thus ensuring that any explored state  
 497 is indeed reachable by a timed run. In doing so, the most challenging part is  
 498 extending the BFS tree wrt a pop. Here, we not only have to find a split of a  
 499 hole into an atomic hole-segment and a hole-segment as in Algorithm 1, but also  
 500 need to keep track of possible partitions of time.

501 **Timed Witness:** As in the untimed case, we generate a witness certifying non-  
 502 emptiness of TMPDA. But, producing a witness for the fix-point computation  
 503 as discussed earlier requires unrolling. The fix-point computation generates a  
 504 pre-computed set **WRT** of tuples  $((s, \nu), t, (s', \nu'))$ , where  $s, s' \in \mathcal{S}$ ,  $t$  is time elapsed  
 505 in the well-nested sequence and  $\nu, \nu' \in \mathbb{N}^{|\mathcal{X}|}$  are integral valuations. This set  
 506 of tuples does not have information about the intermediate transitions and  
 507 time-elapses. To handle this, using the pre-computed information, we define a  
 508 lexicographic progress measure which ensures termination of this search.

509 While the details are in [11] (Algorithm ??), the main idea is as follows:  
 510 the first progress measure is to check if there a time-elapse  $t$  transition possible  
 511 between  $(s, \nu)$  and  $(s', \nu')$  and if so, we print this out. If not,  $\nu' \neq \nu + t$ , and  
 512 some set of clocks have been reset in the transition(s) from  $(s, \nu)$  to  $(s', \nu')$ . The  
 513 second progress measure looks at the sequence of transitions from  $(s, \nu)$  to  $(s', \nu')$ ,  
 514 consisting of reset transitions (at most the number of clocks) that result in  $\nu'$   
 515 from  $\nu$ . If neither the first nor the second progress measure apply, then  $\nu = \nu'$ ,  
 516 and we are left to explore the last progress measure, by exploring at most  $|\mathcal{S}|$   
 517 number of transitions from  $(s, \nu)$  to  $(s', \nu')$ . The lexicographic progress measure  
 518 seamlessly extends the witness generation to the timed setting.

## 519 6 Implementation and Experiments

520 We implemented a tool **BHIM** (**B**ounded **H**oles **I**n **M**PPDA) written in C++ based  
 521 on Algorithm 1, which takes an MPDA and a constant  $K$  as input and returns  
 522 (*True*) iff there exists a  $K$ -hole bounded run from the start state to an accepting  
 523 state of the MPDA. In case there is such an accepting run, **BHIM** generates one  
 524 such, with minimal number of holes. For a given hole bound  $K$ , **BHIM** first tries  
 525 to produce a witness with 0 holes, and iteratively tries to obtain a witness by  
 526 increasing the bound on holes till  $K$ . In most of the cases, **BHIM** found the  
 527 witness before reaching the bound  $K$ . Whenever **BHIM**'s witness had  $K$  holes, it  
 528 is guaranteed that there are no witnesses with a smaller number of holes.

529 To evaluate the performance of **BHIM**, we looked at some available benchmarks  
 530 and modeled them as MPDA. We also added timing constraints to some examples  
 531 such that they can be modeled as TMPDA. Our tests were run on a GNU/Linux  
 532 system with Intel<sup>®</sup> Core<sup>™</sup> i7-4770K CPU @ 3.50GHz, and 16GB of RAM. We  
 533 considered overall 7 benchmarks, of which we sketch 3 in detail here. The details  
 534 of these as well as the remaining ones are in [11].

535 • **Bluetooth Driver** [18]. The Bluetooth device driver example [18], has two  
 536 threads and a shared memory. We model this driver using a 2-stack pushdown  
 537 system, where a state represents the current valuation of the global variables and  
 538 stacks are used to maintain the call-return between different functions and to  
 539 keep the count of processes currently using the driver. There is also a scheduler  
 540 which can preempt any thread executing a non-atomic instruction. A known error  
 541 as pointed out in [18] is a race condition between two threads where one thread  
 542 tries to write to a global variable and the other thread tries to read from it. **BHIM**  
 543 found this error, with a well-nested witness. A timed extension of this example  
 544 was also considered, where, a witness was obtained again with hole bound 0.

545 • **Bluetooth Driver v2** [10,8]. A modified version of Bluetooth driver is  
 546 considered [10,8], where a counter is maintained to count the number of threads  
 547 actively using the driver. A two stack MPDA models this, with one stack simulating  
 548 the counter and another one scheduling the threads. Two known errors reported  
 549 are (i) counter underflow where a counter goes negative, leading to some unwanted  
 550 behavior of the driver, (2) interrupted I/O, where the stopping thread kills the  
 551 driver while the other thread is busy with I/O. The tools SPADE and MAGIC  
 552 [10,8] found one of these two errors, while **BHIM** found both errors, the first using  
 553 a well nested witness, and the second with a 2-hole bounded witness.

554 • **A Multi-threaded Producer Consumer Problem**. The Producer consumer  
 555 problem (see e.g., [25]) is a classic example of concurrency and synchronization.  
 556 A more challenging version of this is when there are multiple producers and  
 557 consumers. Assume that two ingredients called 'A' and 'B' are produced in a  
 558 production line in batches, where a batch can produce arbitrarily many items.  
 559 Further, assume that (1) two units of 'A' and one unit of 'B' make an item called  
 560 'C'; (2) the production line begins its day by producing a batch of A's and then  
 561 the rest of the day, it keeps producing B's in batches, one after the other. During  
 562 the day, 'C's are churned out using 'A' and 'B' in the proportion mentioned

above and, if we run out of ‘A’s, we obtain an error; there is no problem if ‘B’ is exhausted, since a fresh batch producing ‘B’ is commenced.

For  $p, q \in \mathbb{N}$ , consider words of the form  $a^m b^{n_1} (c^2 d)^{n_1} b^{n_2} (c^2 d)^{n_2} \dots b^{n_p} (c^2 d)^{n_p}$  where  $n_i \leq q$  for all  $1 \leq i \leq p$ ,  $a$  represents the production of one unit of ‘A’,  $b$  represents the production of one unit of ‘B’,  $c$  represents consumption of one unit of ‘A’ and  $d$  represents consumption of one unit of ‘B’. Unless  $m \geq 2(n_1 + \dots + n_p)$ , we will obtain an error. This is easily modeled using a 2 stack visibly multi pushdown automaton where  $a, b$  are push symbols of stack 1,2 respectively and  $c, d$  are pop symbols of stack 1,2 respectively. Let  $L_{m,p}$  be the set of words of the above form s.t.  $2(n_1 + \dots + n_{p-1}) < m < 2(n_1 + \dots + n_p)$ ,  $n_i \leq q$ . It can be seen that  $L_{m,p}$  does not have any well-nested word in it, and  $L_{m,p}$  is not scope-bounded. The number of context switches in words of  $L_{m,p}$  depends on the parameters  $m, n_1, \dots, n_p$ . However,  $L_{m,p}$  is 2 hole-bounded : at any position of the word, the open holes come from the unmatched sequences of  $a$  and  $b$  seen so far. BHIM checked for the non-emptiness of  $L_{m,p}$  with a witness of hole bound 2.

- **Critical time constraints [26].** This is one of the timed examples, where we consider the language  $L^{crit} = \{a^y b^z c^y d^z \mid y, z \geq 1\}$  with time constraints between occurrences of symbols. The first  $c$  must appear after 1 time-unit of the last  $a$ , the first  $d$  must appear within 3 time-units of the last  $b$ , and the last  $b$  must appear within 2 time units from the start, and the last  $d$  must appear at 4 time units.  $L^{crit}$  is accepted by a TMPDA with two timed stacks.  $L^{crit}$  has no well-nested word, is 4-context bounded, but only 2 hole-bounded.

- **A Linux Kernel bug dm\_target.c [27].** This example is about a double free bug in the file `drivers/md/dm-target.c` in Linux Kernel 2.5.71, which was introduced to fix a memory leak, but it ended up double freeing the object. BHIM found this bug with a witness of hole bound 3.

- **Concurrent Insertions in Binary Search Trees.** Concurrent insertions in binary search trees is a very important problem in database management systems. [28] proposes an algorithm to solve this problem for concurrent implementations. However, if the locks are not implemented properly, then it is possible for a thread to overwrite others. We modified the algorithm [28] to capture this bug, and modeled it as MPDA. BHIM found the bug with a witness of hole-bound 2.

- **Maze Example.** Finally we consider a robot navigating a maze, picking items; an extended (from single to multiple stack) version of the example from [17]. In the untimed setting, a witness for non-emptiness was obtained with hole-bound 0, while in the extension with time, the witness had a hole-bound 2, since the satisfaction of time constraints required a longer witness.

**Results and Discussion.** The performance of BHIM is presented in Table 1 for untimed examples and in Table 2 for timed examples. Apart from the results in the tables, to check the robustness of BHIM wrt parameters like the number of locations, transitions, stacks, holes and clocks (for TMPDA), we looked at examples with an empty language, by making accepting states non-accepting in the examples considered so far. This forces BHIM to explore all possible paths in the BFS tree, generating the lists at all nodes. The scalability of BHIM wrt all these parameters are in [11].

Name	Locations	Transitions	Stacks	Holes	Time Empty (mili sec)	Time Witness (mili sec)	Memory(KB)
Bluetooth	57	96	2	0	157.9	7.1	7424
Bluetooth v2(err1)	58	99	2	0	27.4	7.1	5096
Bluetooth v2(err2)	58	99	2	2	97.4	24.1	6478
MultiProdCons	11	18	2	2	11.1	0.1	1796
dm-target	13	27	2	3	42.0	5.8	4476
Binary Search Tree	29	78	2	2	60.8	5.1	5143
untimed- $L^{crit}$	6	10	2	2	14.9	0.7	4692
untimed-Maze	9	12	2	0	12.0	0.2	3858
$L^{bh}$ (from Sec. 2.1)	7	13	2	2	22.2	0.6	4404

**Table 1.** Experimental results: Time Empty and Time Witness column represents no. of milliseconds needed for emptiness checking and to generate witness respectively.

Name	Locations	Transitions	Stacks	Clocks	$c_{max}$	Aged(Y/N)	Holes	Time Empty(mili sec)	Time Witness (mili sec)	Memory(KB)
Bluetooth	57	96	2	0	2	Y	0	169.9	101.3	5248
$L^{crit}$	6	10	2	2	8	Y	2	9965.2	3.7	203396
Maze	9	12	2	2	5	Y	2	956.8	9.7	14554

**Table 2.** Experimental results of timed examples. The column  $c_{max}$  is defined as the maximum constant in the automaton, and Aged denotes if the stack is timed or not

**BHIM Vs. State of the art.** What makes BHIM stand apart wrt the existing state of the art tools is that (i) none of the existing tools handle under approximations captured by bounded holes, (ii) none of the existing tools work with multiple stacks in the timed setting (even closed guards!). The state of the art research in underapproximations wrt untimed multistack pushdown systems has produced some amazing tools like GetaFix which handles multi-threaded programs with bounded context switches. While we have adapted some of the examples from GetaFix, the latest available version of GetaFix has some issues in handling those examples<sup>3</sup>. Likewise, SPADE, MAGIC and the counter implementation [27] are currently not maintained. This has come in the way of a performance comparison between BHIM and these tools. Indeed, most examples handled by BHIM correspond to non-context bounded, or non scope bounded, or timed languages which are beyond Getafix. For instance, the 2-hole bounded witness found by BHIM for the language  $L_{20,10}(m = 20, p = 10)$  for the multi producer consumer case cannot be found by GetaFix/MAGIC/SPADE with less than 41 context switches. In the timed setting, the Maze example (TMPDA with 2 clocks, 2 timed stacks) has a 2 hole-bounded witness where the robot visits certain locations an equal number of times. The tool [17] cannot handle this example since it handles only one stack. Lastly, [17] cannot solve binary reachability with an empty stack unlike BHIM.

**BHIM v2.** The next version of BHIM will go symbolic, inspired from GetaFix. The current avatar of BHIM showcases the efficiency of fix-point techniques extended to larger bounded underapproximations; indeed going symbolic will make BHIM much more robust and scalable.

*Acknowledgements.* We would like to thank Gennaro Parlato for the discussions we had on Getafix and for providing us benchmarks.

<sup>3</sup> we did get in touch with the authors, who confirmed this



## References

1. Salvatore La Torre, Parthasarathy Madhusudan, and Gennaro Parlato. A robust class of context-sensitive languages. In *Logic in Computer Science, 2007. LICS 2007. 22nd Annual IEEE Symposium on*, pages 161–170. IEEE, 2007.
2. Salvatore La Torre, Parthasarathy Madhusudan, and Gennaro Parlato. The language theory of bounded context-switching. In *Latin American Symposium on Theoretical Informatics*, pages 96–107. Springer, 2010.
3. Salvatore La Torre, Margherita Napoli, and Gennaro Parlato. Scope-bounded pushdown languages. *International Journal of Foundations of Computer Science*, 27(02):215–233, 2016.
4. Aiswarya Cyriac, Paul Gastin, and K Narayan Kumar. MSO decidability of multi-pushdown systems via split-width. In *International Conference on Concurrency Theory*, pages 547–561. Springer, Berlin, Heidelberg, 2012.
5. Salvatore La Torre and Margherita Napoli. Reachability of multistack pushdown systems with scope-bounded matching relations. In *International Conference on Concurrency Theory*, page 203–218. Springer, 2011.
6. Salvatore La Torre and Parlato Gennaro. Scope-bounded multistack pushdown systems: Fixed-point, sequentialization, and tree-width. 2012.
7. Salvatore La Torre, Madhusudan Parthasarathy, and Gennaro Parlato. Analyzing recursive programs using a fixed-point calculus. *ACM Sigplan Notices*, 44(6):211–222, 2009.
8. Gaël Patin, Mihaela Sighireanu, and Tayssir Touili. Spade: Verification of multithreaded dynamic and recursive programs. In *International Conference on Computer Aided Verification*, pages 254–257. Springer, 2007.
9. Shaz Qadeer. The case for context-bounded verification of concurrent programs. In *Model Checking Software, 15th International SPIN Workshop, Los Angeles, CA, USA, August 10-12, 2008, Proceedings*, pages 3–6, 2008.
10. Sagar Chaki, Edmund Clarke, Nicholas Kidd, Thomas Reps, and Tayssir Touili. Verifying concurrent message-passing C programs with recursive calls. In *International Conference on Tools and Algorithms for the Construction and Analysis of Systems*, page 334–349. Springer, 2006.
11. Akshay S, Gastin Paul, S Krishna, and Roychowdhury Sparsa. Supplementary material: Revisiting under-approximate reachability in MPDA. Available at <https://cse.iitb.ac.in/~sparsa/bhim/>, 2019.
12. Mohamed Faouzi Atig. Model-checking of ordered multi-pushdown automata. *arXiv preprint arXiv:1209.1916*, 2012.
13. Ahmed Bouajjani, Rachid Echahed, and Riadh Robbana. On the automatic verification of systems with continuous variables and unbounded discrete data structures. In *International Hybrid Systems Workshop*, pages 64–85. Springer, 1994.
14. Parosh Aziz Abdulla, Mohamed Faouzi Atig, and Jari Stenman. Dense-timed pushdown automata. In *Proceedings of the 27th Annual IEEE Symposium on Logic in Computer Science, LICS 2012, Dubrovnik, Croatia, June 25-28, 2012*, page 35–44, 2012.
15. Parosh Aziz Abdulla, Mohamed Faouzi Atig, and Jari Stenman. The minimal cost reachability problem in priced timed pushdown systems. In *Language and Automata Theory and Applications - 6th International Conference, LATA 2012, A Coruña, Spain, March 5-9, 2012. Proceedings*, pages 58–69, 2012.
16. S. Akshay, Paul Gastin, and Shankara Narayanan Krishna. Analyzing Timed Systems Using Tree Automata. *Logical Methods in Computer Science*, Volume 14, Issue 2, May 2018.

- 684 17. S. Akshay, Paul Gastin, Shankara Narayanan Krishna, and Ilias Sarkar. Towards  
685 an efficient tree automata based technique for timed systems. In *28th International  
686 Conference on Concurrency Theory, CONCUR 2017, September 5-8, 2017, Berlin,  
687 Germany*, pages 39:1–39:15, 2017.
- 688 18. Shaz Qadeer and Dinghao Wu. Kiss: keep it simple and sequential. *Acm sigplan  
689 notices*, 39(6):14–24, 2004.
- 690 19. P Madhusudan and Gennaro Parlato. The tree width of auxiliary storage. In *ACM  
691 SIGPLAN Notices*, volume 46, pages 283–294. ACM, 2011.
- 692 20. S. Akshay, Paul Gastin, Vincent Jugé, and Shankara Narayanan Krishna. Timed  
693 systems through the lens of logic. In *34th Annual ACM/IEEE Symposium on Logic  
694 in Computer Science, LICS 2019, Vancouver, BC, Canada, June 24-27, 2019*, pages  
695 1–13, 2019.
- 696 21. Aiswarya Cyriac. *Verification of communicating recursive programs via split-width.  
697 (Vérification de programmes récurifs et communicants via split-width)*. PhD thesis,  
698 École normale supérieure de Cachan, France, 2014.
- 699 22. Rajeev Alur and Parthasarathy Madhusudan. Visibly pushdown languages. In  
700 *Proceedings of the thirty-sixth annual ACM symposium on Theory of computing*,  
701 pages 202–211. ACM, 2004.
- 702 23. Zhe Dang, Oscar H Ibarra, Tevfik Bultan, Richard A Kemmerer, and Jianwen Su.  
703 Binary reachability analysis of discrete pushdown timed automata. In *International  
704 Conference on Computer Aided Verification*, page 69–84. Springer, 2000.
- 705 24. Thomas H Cormen, Charles E Leiserson, Ronald L Rivest, and Clifford Stein.  
706 *Introduction to algorithms*. MIT press, 2009.
- 707 25. Abraham Silberschatz, Greg Gagne, and Peter B Galvin. *Operating system concepts*.  
708 Wiley, 2018.
- 709 26. Devendra Bhave, Vrunda Dave, Shankara Narayanan Krishna, Ramchandra  
710 Phawade, and Ashutosh Trivedi. A perfect class of context-sensitive timed languages.  
711 In *International Conference on Developments in Language Theory*, pages 38–50.  
712 Springer, Berlin, Heidelberg, 2016.
- 713 27. Matthew Hague and Anthony Widjaja Lin. Synchronisation- and reversal-bounded  
714 analysis of multithreaded programs with counters. In *Computer Aided Verification  
715 - 24th International Conference, CAV 2012, Berkeley, CA, USA, July 7-13, 2012  
716 Proceedings*, page 260–276, 2012.
- 717 28. HT Kung and Philip L Lehman. Concurrent manipulation of binary search trees.  
718 *ACM Transactions on Database Systems (TODS)*, 5(3):354–382, 1980.